



Attention Models in Deep Learning



By
Chandresh Kumar Maurya
CSE deptt, IIT Indore

<https://chandu8542.github.io/>




Slides credit

- 1. Prof. Ming Li, Deep Learning and NLP course, UoW**
- 2. Amaia Salvador, DCU, Dublin**
- 3. Prof. Chris Manning, Stanford NLP course CS 224n**

CONTENT

1. Introduction and Motivation
2. Seq-2-seq
3. Attention
4. Transformers


1. Fully connected network, feedforward network



To learn the weights on the edges

2. CNN

A CNN is a neural network with some convolutional layers (and some other layers). A convolutional layer has a number of filters that do convolutional operation.



Convolutional layer

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

Input

These are the network parameters to be learned.

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2

⋮ ⋮

Each filter detects a small pattern (3 x 3).

Convolution Operation


stride=1

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

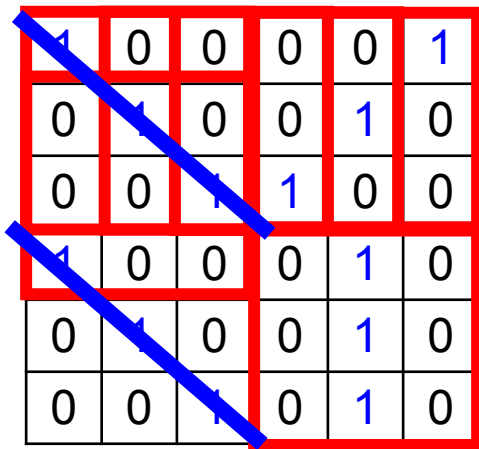
Dot
produc
t



Input

Convolution

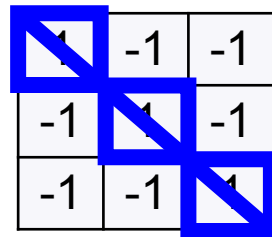
stride=1



A 6x6 grid of numbers. A 3x3 red border highlights the top-left corner. A blue diagonal line runs from the top-left to the bottom-right.

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

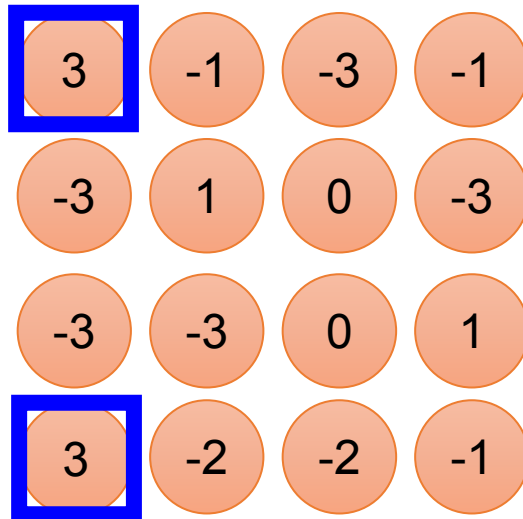
Input



A 3x3 grid of numbers. A blue diagonal line runs from the top-left to the bottom-right.

1	-1	-1
-1	1	-1
-1	-1	1


Filter 1



A 4x4 grid of orange circles containing numbers. The top-left and bottom-left circles are highlighted with blue squares.


3	-1	-3	-1
-3	1	0	-3
-3	-3	0	1
3	-2	-2	-1

3. RNN




Parameters to be learned:
 U, V, W

Simple RNN vs LSTM



(a) RNN




(b) LSTM

Attention Models: Motivation




Image:
 $H \times W \times 3$



Attention Models: Motivation



Image:
 $H \times W \times 3$




What is wrong here?

Attention Models: Motivation




Image:
 $H \times W \times 3$



What is wrong here? The whole input volume is used to predict the output... despite the fact that not all pixels are equally important


Attention Models: Motivation



Attention models can
relieve computational burden

Helpful when processing big
images !

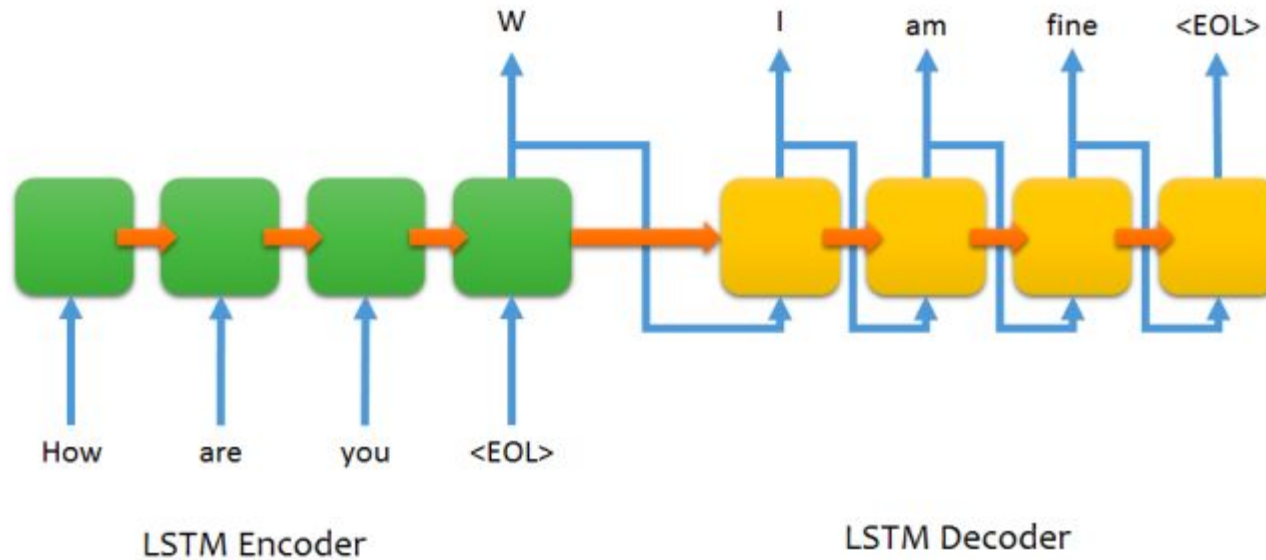
Attention Models: Motivation



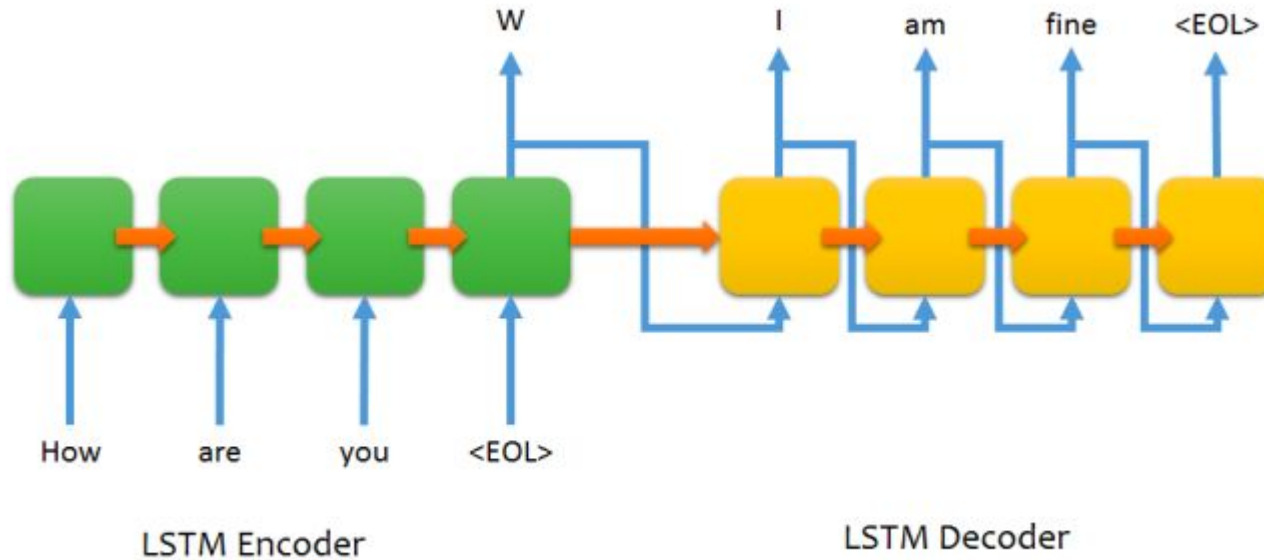
Attention models can
relieve computational burden

Helpful when processing big
images !

Encoder-Decoder LSTM structure for chatting (for non-intelligent beings)

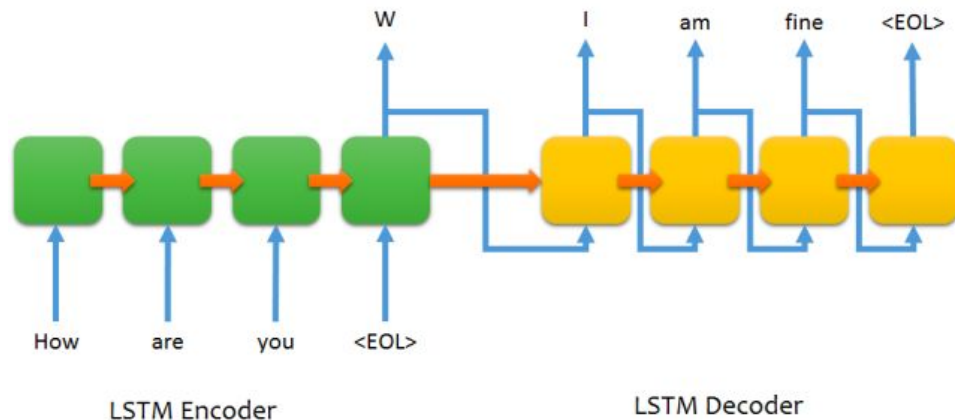


Encoder-Decoder LSTM structure for chatting (for non-intelligent beings)



What is wrong here?

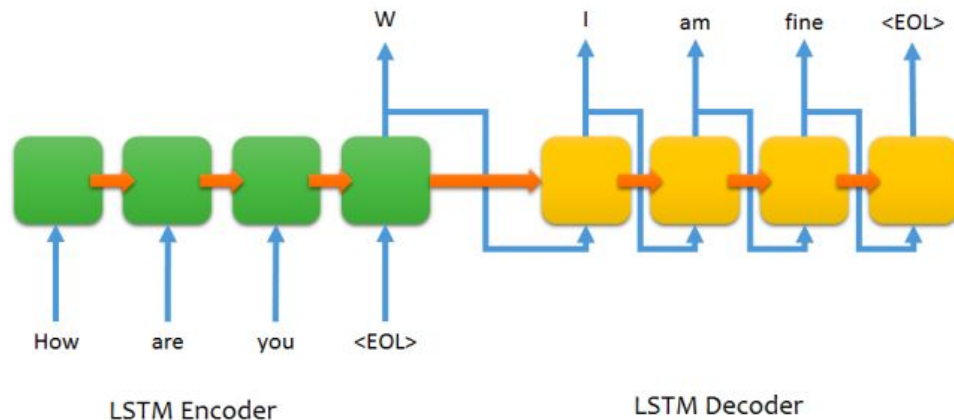
Encoder-Decoder LSTM structure for chatting (for non-intelligent beings)



What is wrong here?

Limitation 1: The whole information is encoded in a fixed-size vector, no matter the length of the input sentence.

Encoder-Decoder LSTM structure for chatting (for non-intelligent beings)



What is wrong here?

Limitation 2: All output predictions are based on the final and static recurrent state of the encoder (h_T). No matter the output word being predicted at each time step, all input words are considered in an equal way.


Attention!

- **Attention** provides a solution to the bottleneck problem.
- Core idea: on each step of the decoder, *use direct connection to the encoder to focus on a particular part* of the source sequence




- First, we will show via diagram (no equations), then we will show with equations


Seq-2-seq with Attention




Seq-2-seq with Attention



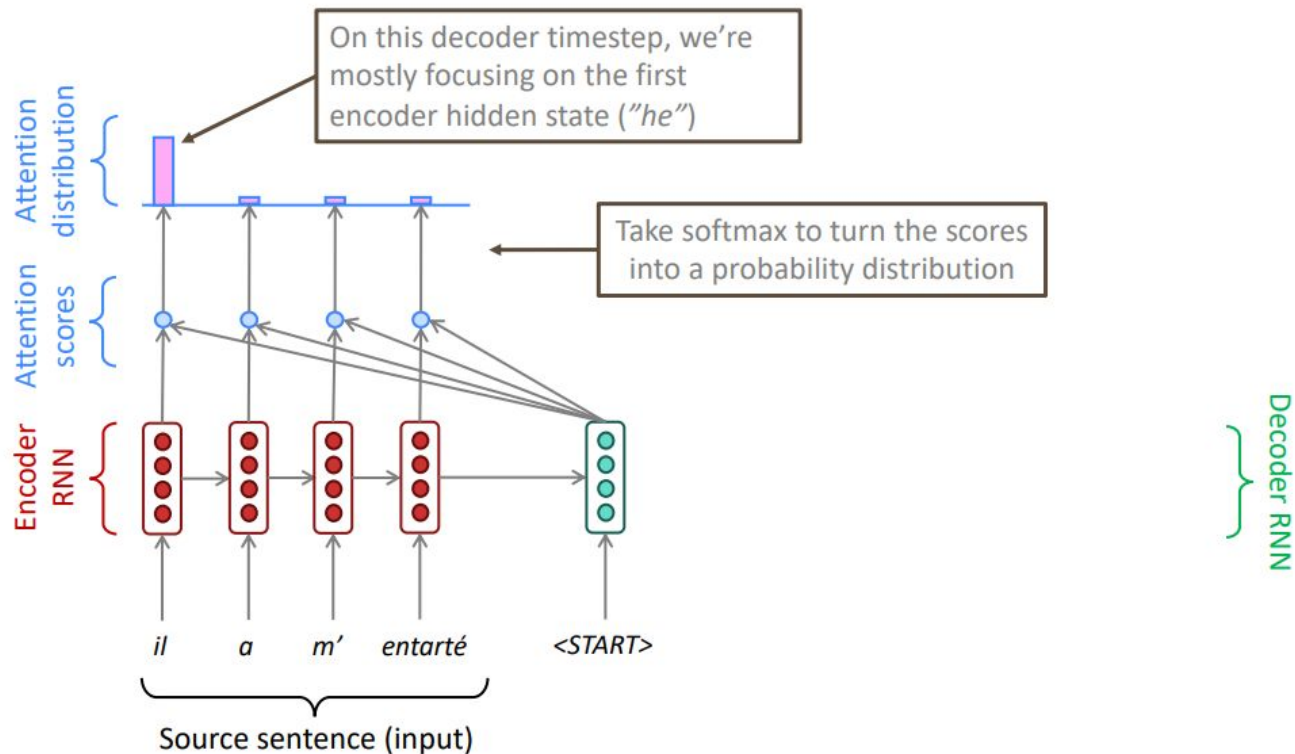
Seq-2-seq with Attention



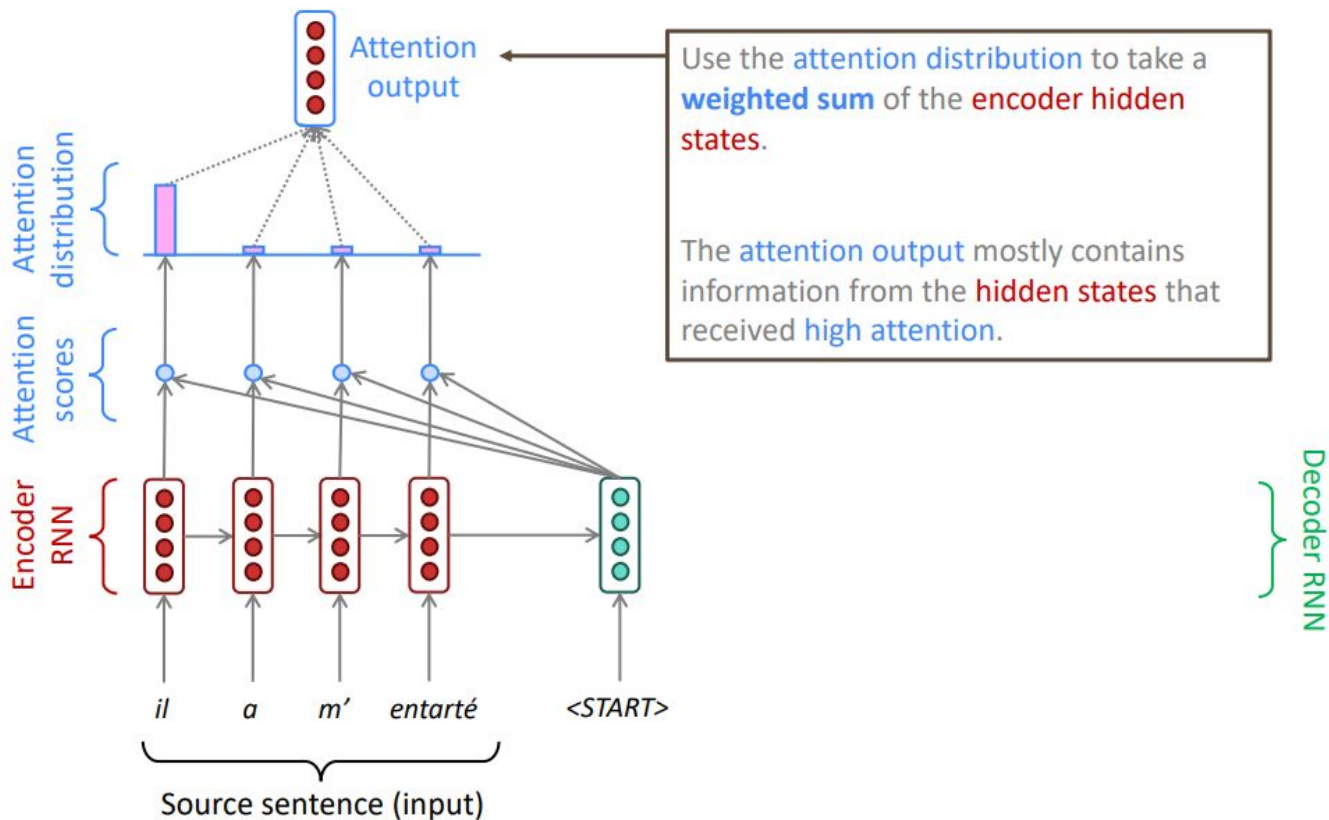
Seq-2-seq with Attention



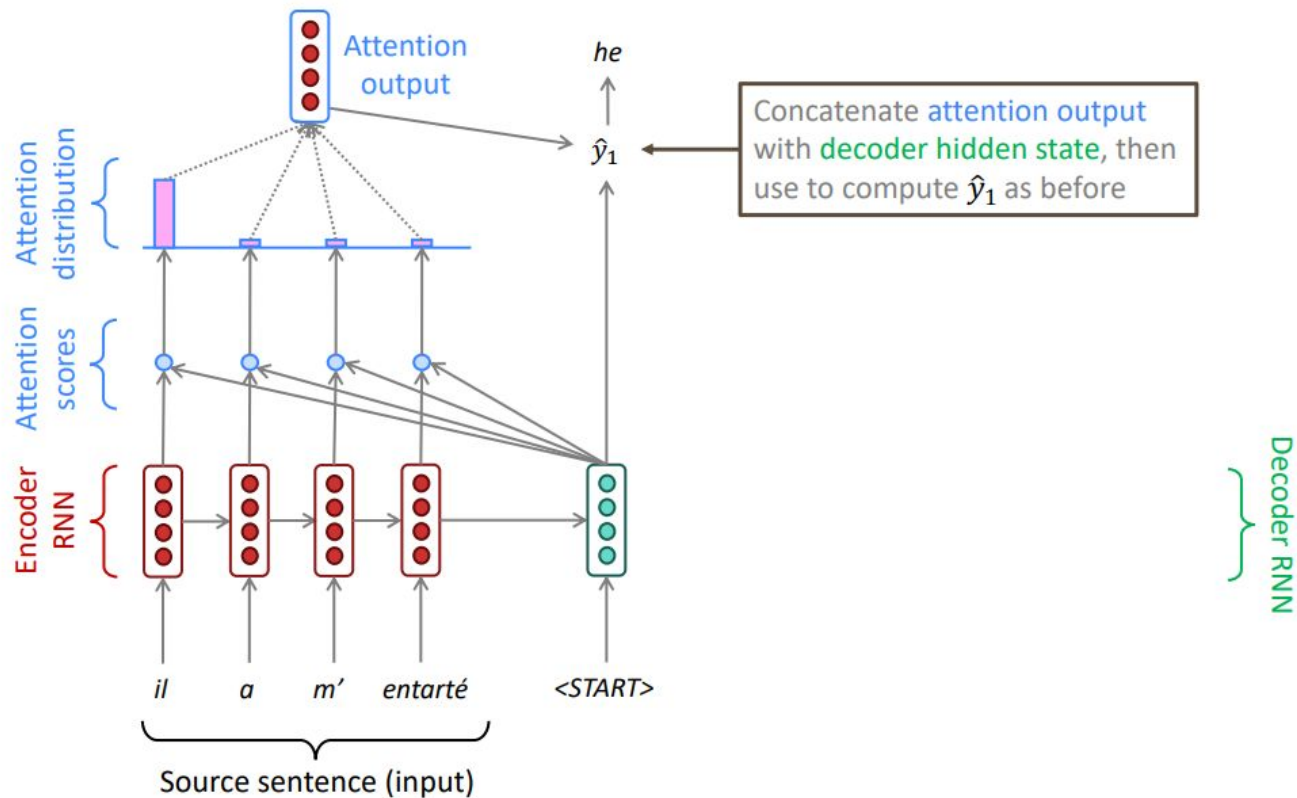
Seq-2-seq with Attention




Seq-2-seq with Attention



Seq-2-seq with Attention

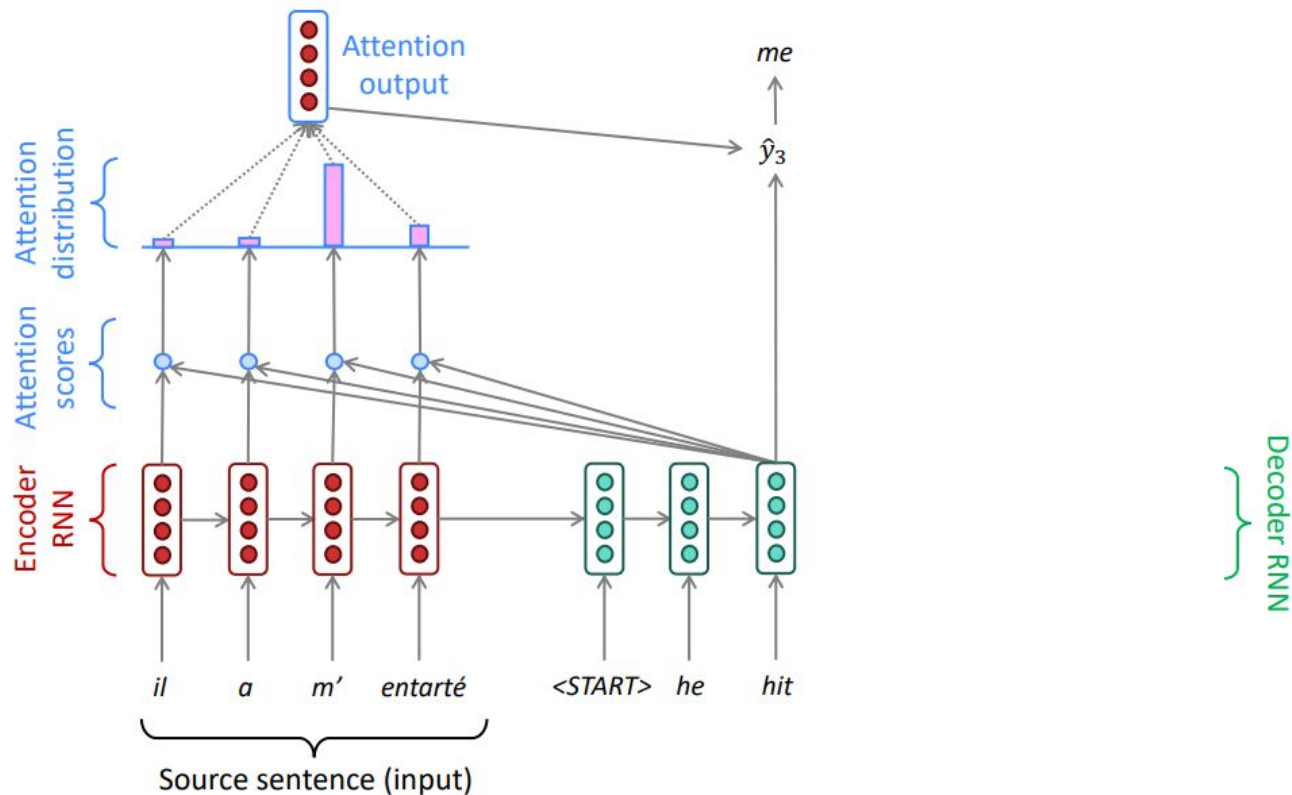


Seq-2-seq with Attention

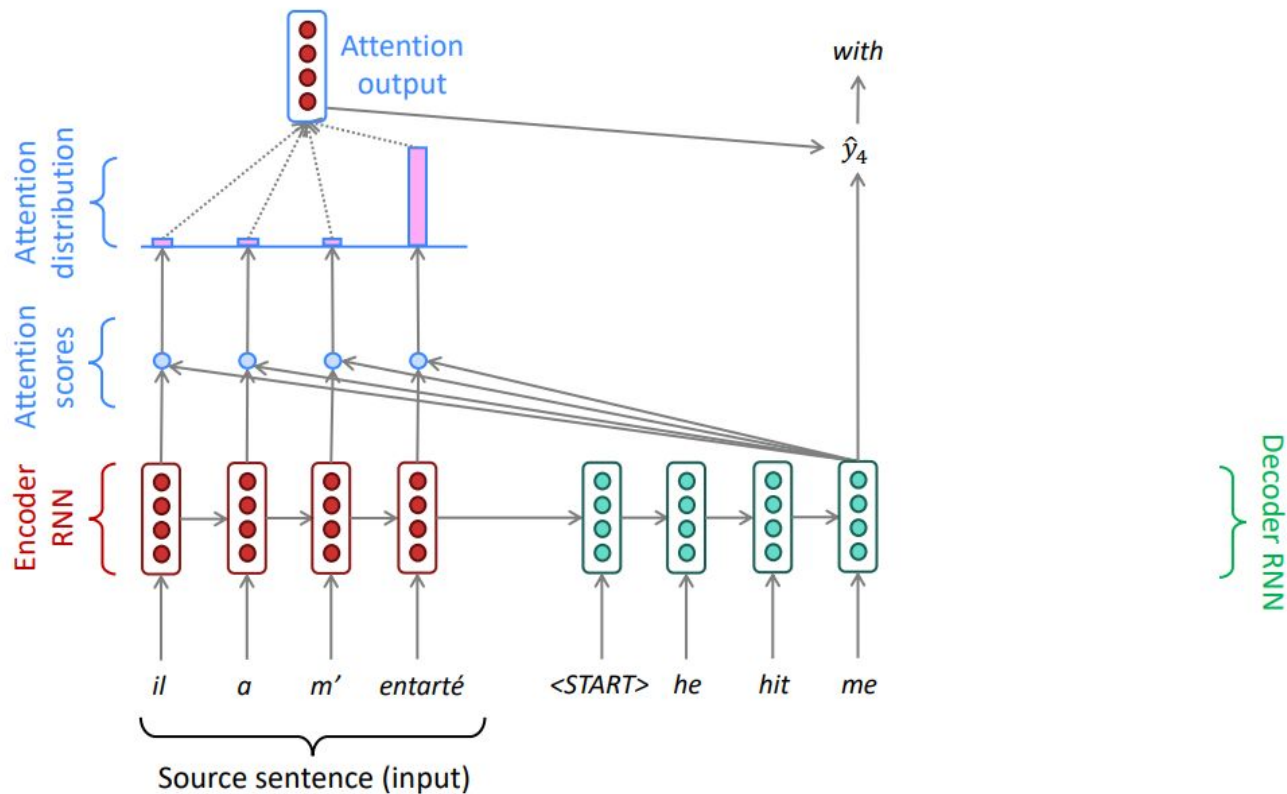


Sometimes we take the attention output from the previous step, and also feed it into the decoder (along with the usual decoder input).

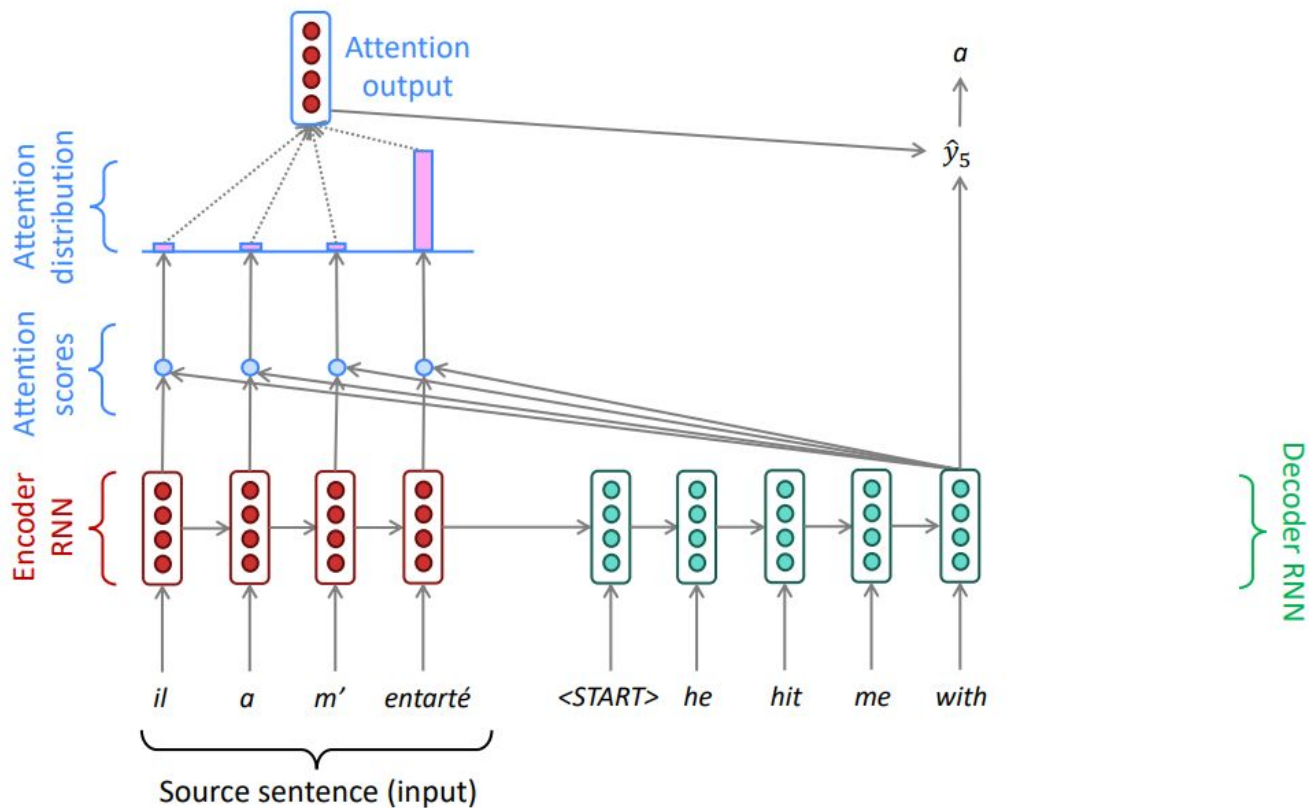
Seq-2-seq with Attention



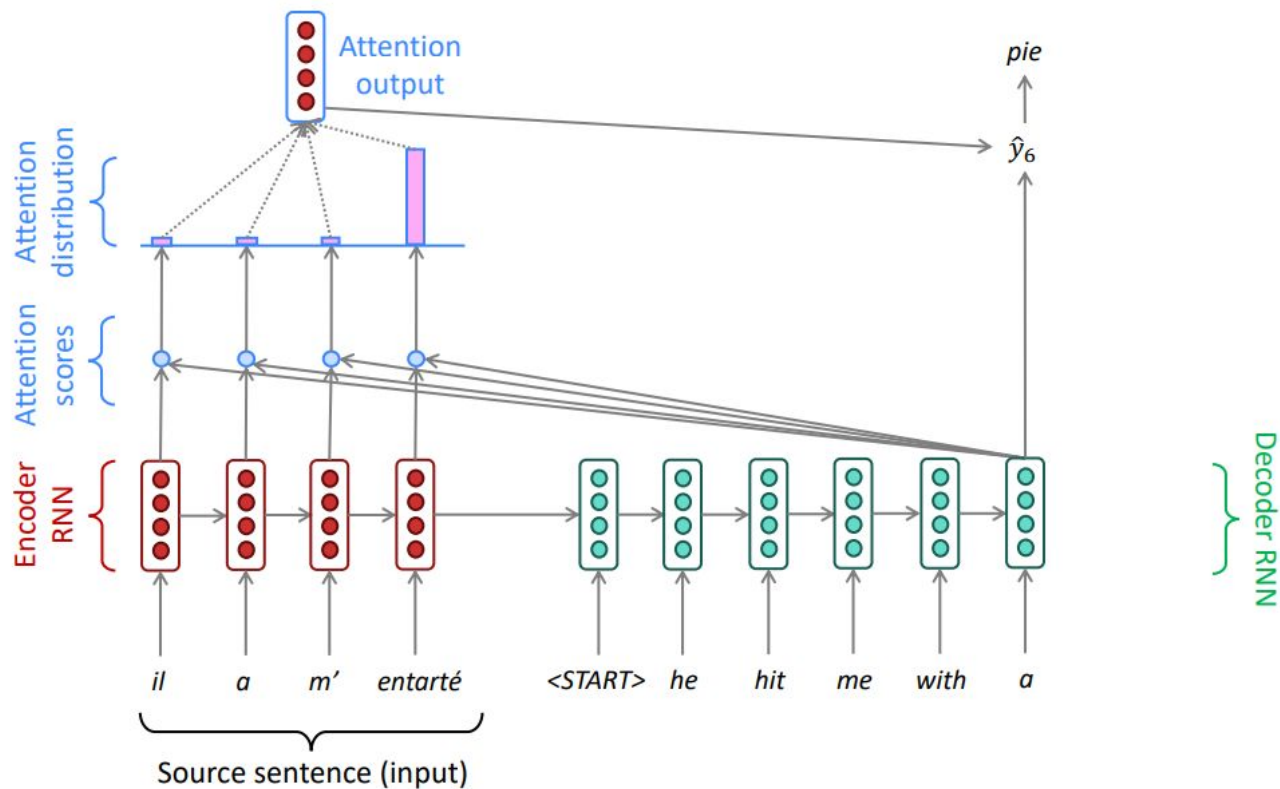
Seq-2-seq with Attention



Seq-2-seq with Attention



Seq-2-seq with Attention



Attention: in equations

- We have encoder hidden states $h_1, \dots, h_N \in \mathbb{R}^h$
- On timestep t , we have decoder hidden state $s_t \in \mathbb{R}^h$
- We get the attention scores e^t for this step:

$$e^t = [s_t^T h_1, \dots, s_t^T h_N] \in \mathbb{R}^N$$

- We take softmax to get the attention distribution α^t for this step (this is a probability distribution and sums to 1)

$$\alpha^t = \text{softmax}(e^t) \in \mathbb{R}^N$$

- We use α^t to take a weighted sum of the encoder hidden states to get the attention output a_t

$$a_t = \sum_{i=1}^N \alpha_i^t h_i \in \mathbb{R}^h$$

- Finally we concatenate the attention output a_t with the decoder hidden state s_t and proceed as in the non-attention seq2seq model

$$[a_t; s_t] \in \mathbb{R}^{2h}$$

Attention is Great

- Attention significantly **improves NMT performance**
 - It's very useful to allow decoder to focus on certain parts of the source
- Attention **solves the bottleneck problem**
 - Attention allows decoder to look directly at source; bypass bottleneck
- Attention **helps with vanishing gradient problem**
 - Provides shortcut to faraway states
- Attention provides **some interpretability**
 - By inspecting attention distribution, we can see what the decoder was focusing on
 - We get (soft) **alignment for free!**
 - This is cool because we never explicitly trained an alignment system
 - The network just learned alignment by itself

	he	hit	me	with	a	pie
il						
a						
m'						
entarté						

Attention is a general Deep Learning technique

- We've seen that attention is a great way to improve the sequence-to-sequence model for Machine Translation.
 - However: You can use attention in many architectures (not just seq2seq) and many tasks (not just MT)
- More general definition of attention:
 - Given a set of vector *values*, and a vector *query*, attention is a technique to compute a weighted sum of the values, dependent on the query.
- We sometimes say that the *query attends to the values*.
 - For example, in the seq2seq + attention model, each decoder hidden state (query) *attends to* all the encoder hidden states (values).

Attention is a general Deep Learning technique

More general definition of attention:

Given a set of vector *values*, and a vector *query*, attention is a technique to compute a weighted sum of the values, dependent on the query.

Intuition:

- The weighted sum is a *selective summary* of the information contained in the values, where the query determines which values to focus on.
- Attention is a way to obtain a *fixed-size representation of an arbitrary set of representations* (the values), dependent on some other representation (the query).

There are several attention variants

- We have some *values* $\mathbf{h}_1, \dots, \mathbf{h}_N \in \mathbb{R}^{d_1}$ and a *query* $\mathbf{s} \in \mathbb{R}^{d_2}$

- Attention always involves:

1. Computing the *attention scores*

$$\mathbf{e} \in \mathbb{R}^N$$

There are
multiple ways
to do this

2. Taking softmax to get *attention distribution* α :

$$\alpha = \text{softmax}(\mathbf{e}) \in \mathbb{R}^N$$

3. Using attention distribution to take weighted sum of values:

$$\mathbf{a} = \sum_{i=1}^N \alpha_i \mathbf{h}_i \in \mathbb{R}^{d_1}$$

thus obtaining the *attention output* \mathbf{a} (sometimes called the *context vector*)

There are several attention variants

There are several ways you can compute $e \in \mathbb{R}^N$ from $\mathbf{h}_1, \dots, \mathbf{h}_N \in \mathbb{R}^{d_1}$ and $\mathbf{s} \in \mathbb{R}^{d_2}$:

- Basic dot-product attention: $e_i = \mathbf{s}^T \mathbf{h}_i \in \mathbb{R}$
 - Note: this assumes $d_1 = d_2$
 - This is the version we saw earlier
- Multiplicative attention: $e_i = \mathbf{s}^T \mathbf{W} \mathbf{h}_i \in \mathbb{R}$
 - Where $\mathbf{W} \in \mathbb{R}^{d_2 \times d_1}$ is a weight matrix
- Additive attention: $e_i = \mathbf{v}^T \tanh(\mathbf{W}_1 \mathbf{h}_i + \mathbf{W}_2 \mathbf{s}) \in \mathbb{R}$
 - Where $\mathbf{W}_1 \in \mathbb{R}^{d_3 \times d_1}$, $\mathbf{W}_2 \in \mathbb{R}^{d_3 \times d_2}$ are weight matrices and $\mathbf{v} \in \mathbb{R}^{d_3}$ is a weight vector.
 - d_3 (the attention dimensionality) is a hyperparameter

Image caption generation using attention




Image caption generation using attention







Image caption generation using attention



Compute alignments scores (scalars):

$$e_{t,i,j} = f_{att}(h_{t-1}, z_{i,j})$$

$f_{att}(\cdot)$ is an MLP

Normalize to get attention weights:

$$a_{t,i,j} = \text{softmax}(e_{t,i,j})$$

$0 < a_{t,i,j} < 1$,
attention values sum to 1

Image caption generation using attention




A woman is throwing a frisbee in a park.



A dog is standing on a hardwood floor.



A stop sign is on a road with a mountain in the background.



A little girl sitting on a bed with a teddy bear.



A group of people sitting on a boat in the water.



A giraffe standing in a forest with trees in the background.

Kelvin Xu, Jimmy Ba, Ryan Kiros, Kyunghyun Cho, Aaron Courville, Ruslan Salakhutdinov, Richard Zemel, Yoshua Bengio,

“Show, Attend and Tell: Neural Image Caption Generation with Visual Attention”, ICML, 2015

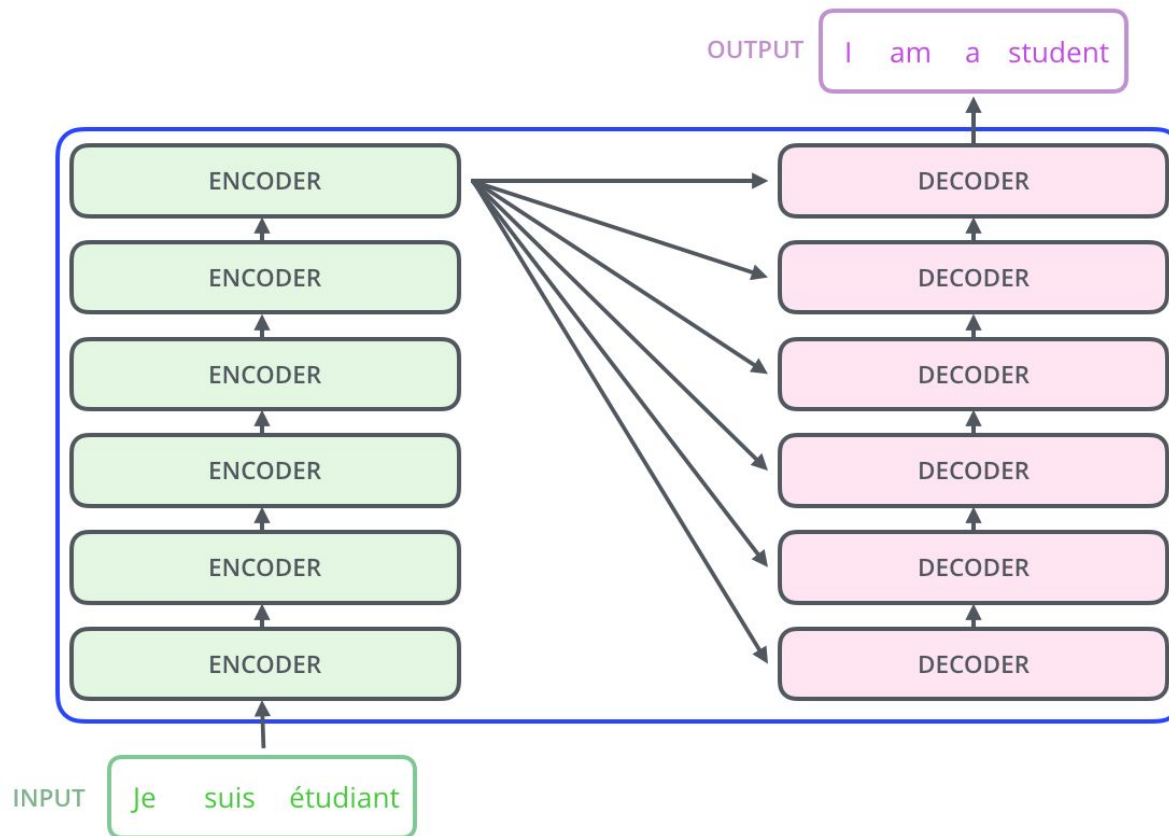
Attention and Transformers

Many new ideas

1. ULM-FiT, pre-training, transfer learning in NLP
2. Recurrent models require linear sequential computation, hard to parallelize. ELMo, bidirectional LSTM.
3. In order to reduce such sequential computation, several models based on CNN are introduced, such as ConvS2S and ByteNet. Dependency for ConvS2S needs linear depth, and ByteNet logarithmic.
4. The transformer is the first transduction model relying entirely on self-attention to compute the representations of its input and output without using RNN or CNN.

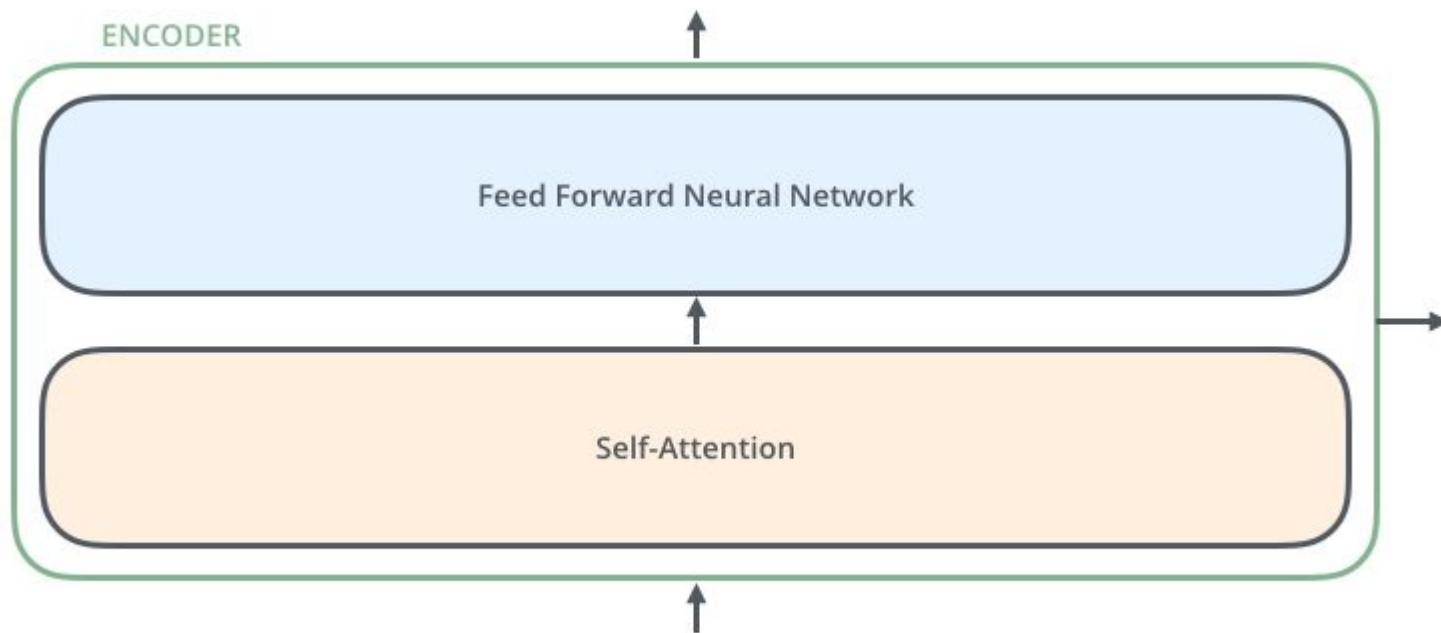
Attention and Transformers

Transformer



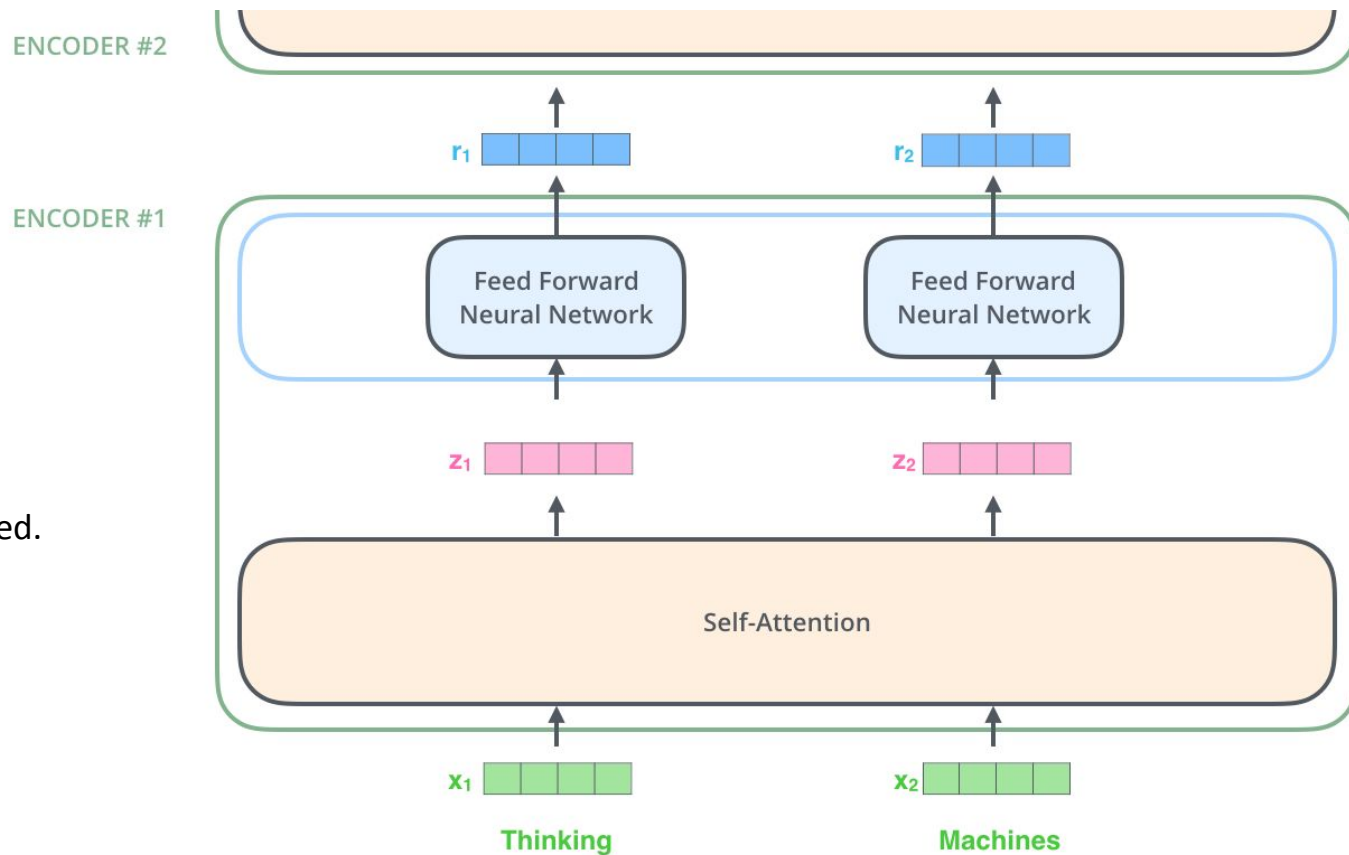
Attention and Transformers

An Encoder Block: same structure, different parameters



Attention and Transformers

Encoder



Note: The ffnn is independent for each word. Hence can be parallelized.

Attention and Transformers

Self

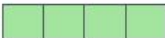
Attention

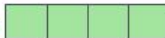
Input

Thinking

Machines

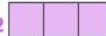
Embedding


X_1 

X_2 

Queries

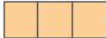
q_1 

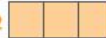
q_2 




W^Q

Keys

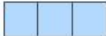
k_1 

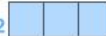
k_2 




W^K

Values

v_1 

v_2 



W^V

First we create three vectors
by multiplying input
embedding
(1x512)
 x_i with three matrices
(64x512):

$$q_i = x_i W^Q$$

$$k_i = x_i W^K$$

$$v_i = x_i W^V$$

Attention and Transformers

Self

Attention

Now we need to calculate a score to determine how much focus to place on other Parts of the input.

Input

Embedding


Queries

Keys

Values


Score

Thinking



$$q_1 \cdot k_1 = 112$$

Machines



$$q_1 \cdot k_2 = 96$$

Attention and Transformers

Self Attention

Formula

$$\text{softmax}\left(\frac{Q \times K^T}{\sqrt{d_k}}\right) V$$
$$= Z$$

$d_k=64$ is dimension of key vector

Input

Embedding

Queries

Keys

Values

Score

Divide by 8 ($\sqrt{d_k}$)

Softmax

Softmax

X

Value

Sum

Thinking

Machines

x_1

x_2

q_1

q_2

k_1

k_2

v_1

v_2

$q_1 \cdot k_1 = 112$

$q_1 \cdot k_2 = 96$

14

12

0.88

0.12

\tilde{v}_1

\tilde{v}_2

$z_1 = 0.88v_1 + 0.12v_2$

z_1

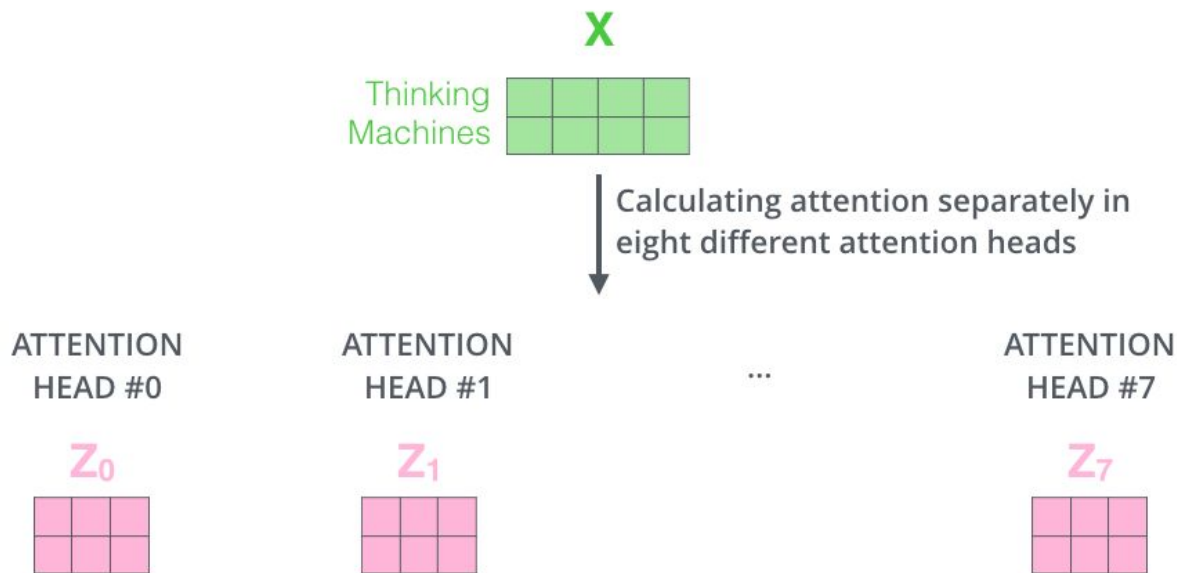
z_2

Attention and Transformers

Multiple

heads

1. It expands the model's ability to focus on different positions.
2. It gives the attention layer multiple "representation subspaces"



Attention and Transformers

1) Concatenate all the attention heads




3) The result would be the Z matrix that captures information from all the attention heads. We can send this forward to the FFNN



2) Multiply with a weight matrix W^O that was trained jointly with the model

\times

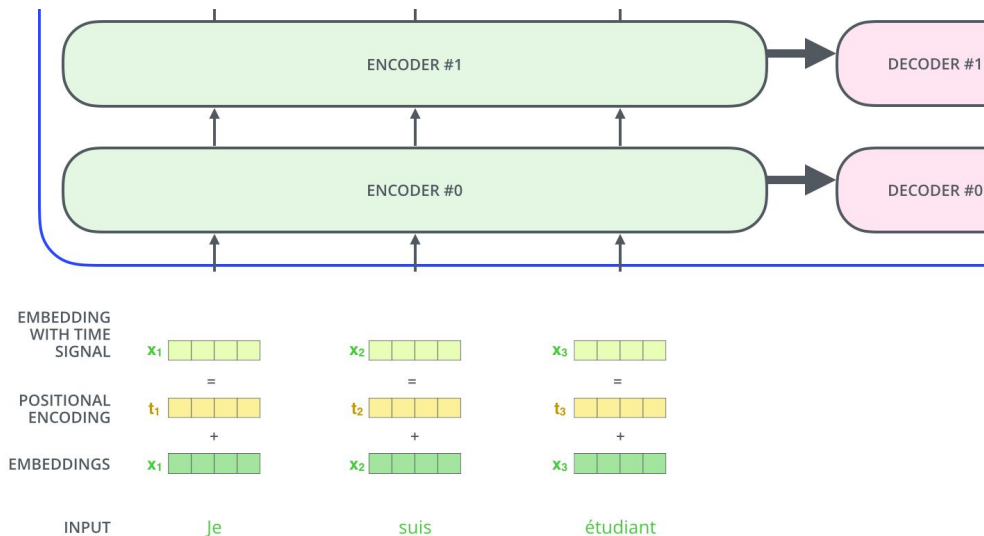


Attention and Transformers

Representing the input order (positional encoding)

The transformer adds a vector to each input embedding. These vectors follow a specific pattern that the model learns, which helps it determine the position of each word, or the distance between different words in the sequence. The intuition here is that adding these values to the embeddings provides meaningful distances between the embedding vectors once they're projected into Q/K/V vectors and during dot-product attention.

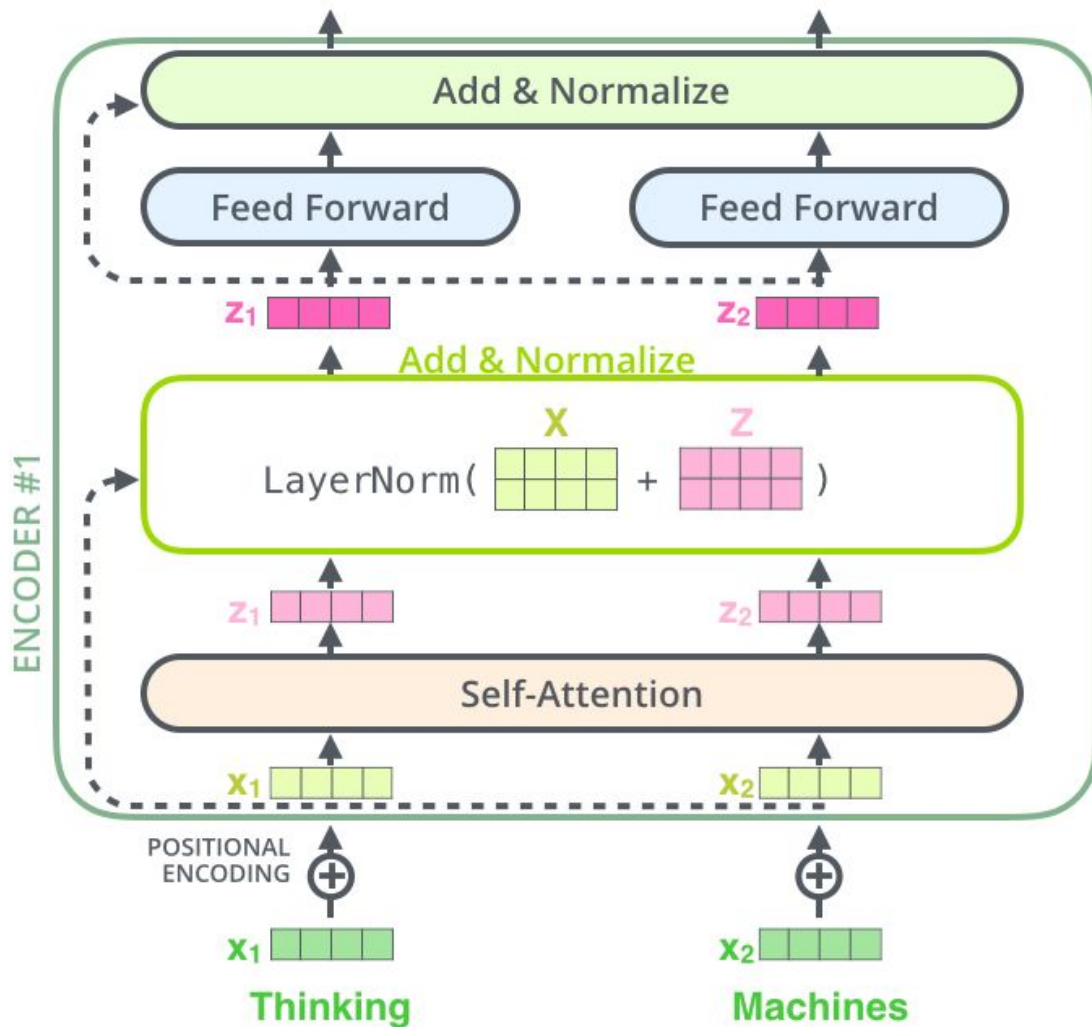
Can somebody present positional encoding following
https://kazemnejad.com/blog/transformer_architecture_positional_encoding/



Attention and Transformers

Add and Normalize

In order to regulate the computation, this is a normalization layer so that each feature (column) have the same average and deviation.




Attention and Transformers

Layer Normalization

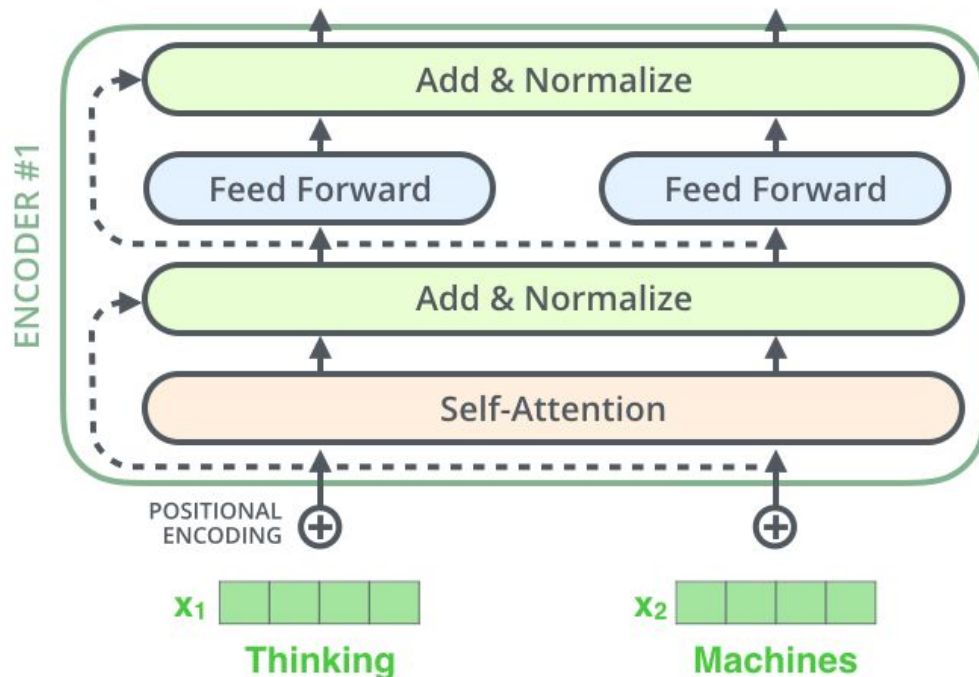

(Hinton)

Layer normalization normalizes the inputs across the features.

Batch Normalization

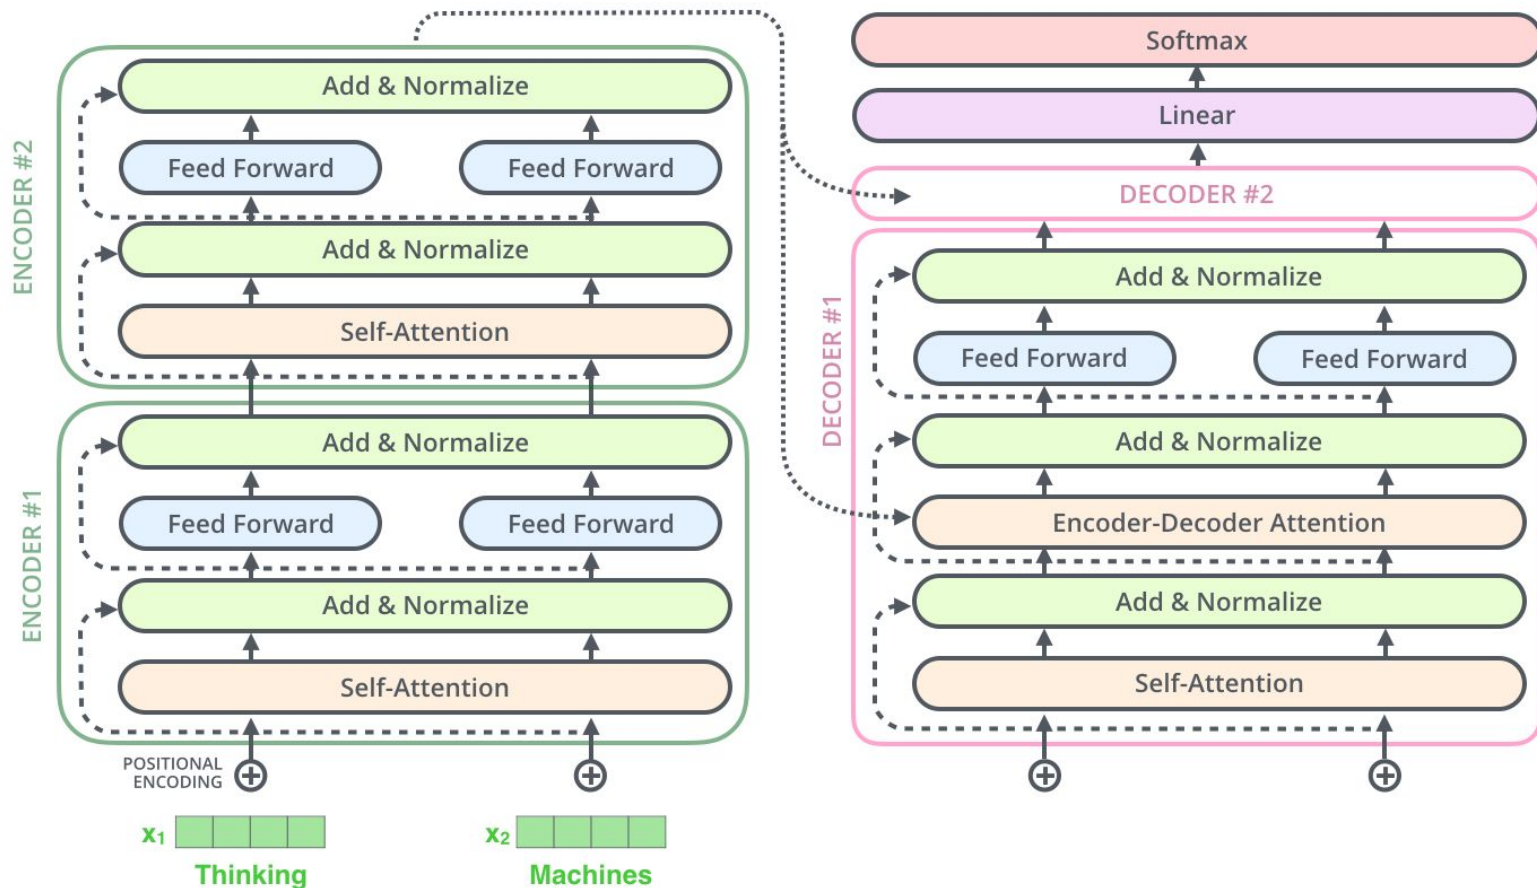


Layer Normalization



The complete transformer

The encoder-decoder attention is just like self attention, except it uses K, V from the top of encoder output, and its own Q



Attention and Transformers

Which word in our vocabulary
is associated with this index?

Decoder's

Output

Linear


Layer

Get the index of the cell
with the highest value
(argmax)

am


5

log_probs



Softmax

logits



Linear

Decoder stack output

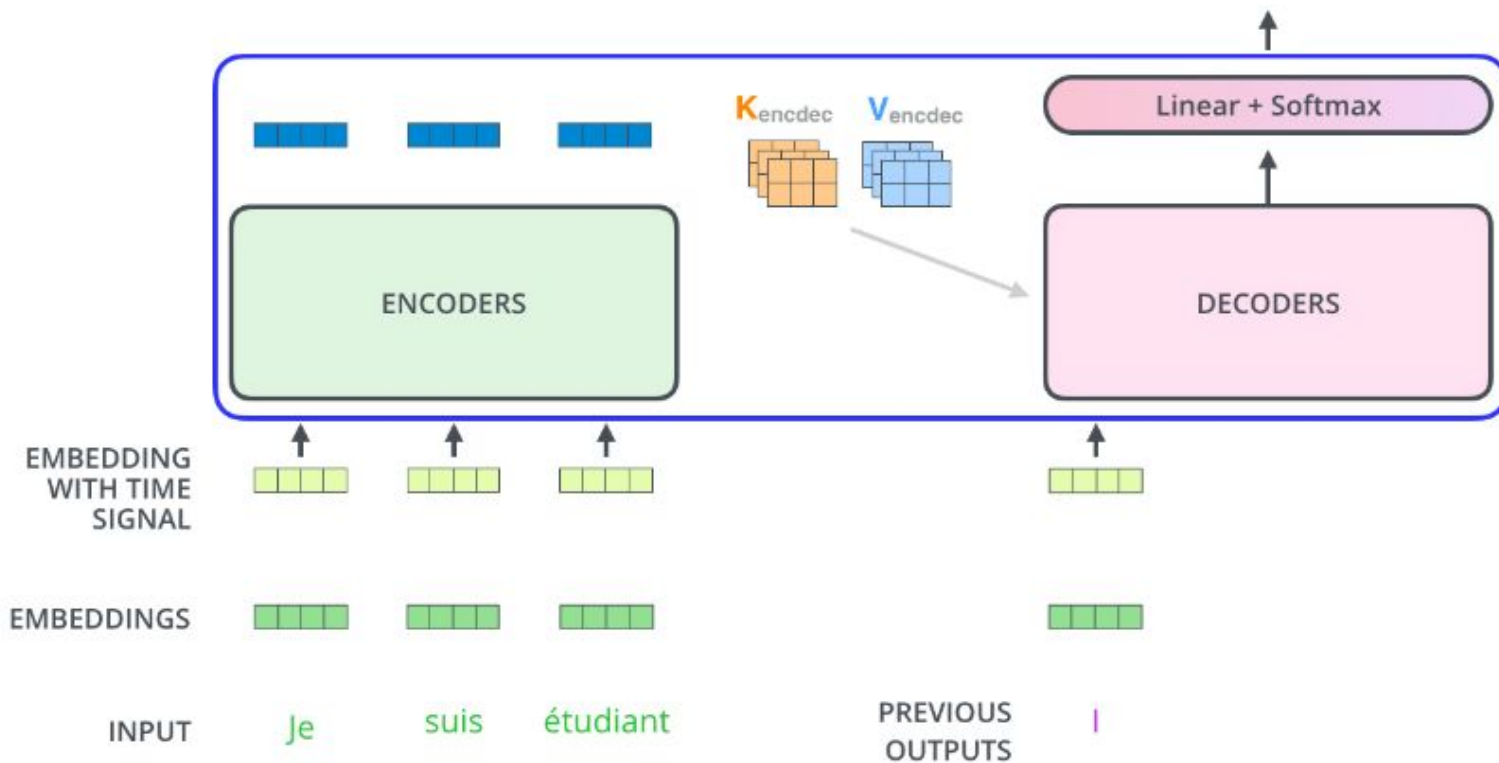


Attention and Transformers

How it works

Decoding time step: 1 2 3 4 5 6

OUTPUT | am



But what about Self-attention?

Attention and Transformers

But what about self-attention when the input is
“incomplete” ?

The solution is to set future unknown values with
“ $-\infty$ ” .

The same for Encoder-Decoder Attention.

Attention and Transformers

Training and the Loss Function

Untrained Model Output



Correct and desired output



a


am

I

thanks

student

<eos>



We can use cross Entropy.

We can also optimize two words at a time: using BEAM search: keep a few alternatives for the first word.

Attention and Transformers

Cross Entropy and KL (Kullback-Leibler) divergence

- **Entropy**: $E(P) = - \sum_i P(i) \log P(i)$ - expected code length (also optimal)
- **Cross Entropy**: $C(P) = - \sum_i P(i) \log Q(i)$ – expected coding length using optimal code for Q
- **KL divergence**:
 $D_{KL}(P \parallel Q) = \sum_i P(i) \log[P(i)/Q(i)] = \sum_i P(i) [\log P(i) - \log Q(i)]$, extra bits
- **JSD** $(P \parallel Q) = \frac{1}{2} D_{KL}(P \parallel M) + \frac{1}{2} D_{KL}(Q \parallel M)$, $M = \frac{1}{2} (P + Q)$, symmetric KL

* JSD = Jensen-Shannon Divergency

Attention and Transformers

Transformer Results

Table 2: The Transformer achieves better BLEU scores than previous state-of-the-art models on the English-to-German and English-to-French newstest2014 tests at a fraction of the training cost.

Model	BLEU		Training Cost (FLOPs)	
	EN-DE	EN-FR	EN-DE	EN-FR
ByteNet [18]	23.75			
Deep-Att + PosUnk [39]		39.2		$1.0 \cdot 10^{20}$
GNMT + RL [38]	24.6	39.92	$2.3 \cdot 10^{19}$	$1.4 \cdot 10^{20}$
ConvS2S [9]	25.16	40.46	$9.6 \cdot 10^{18}$	$1.5 \cdot 10^{20}$
MoE [32]	26.03	40.56	$2.0 \cdot 10^{19}$	$1.2 \cdot 10^{20}$
Deep-Att + PosUnk Ensemble [39]		40.4		$8.0 \cdot 10^{20}$
GNMT + RL Ensemble [38]	26.30	41.16	$1.8 \cdot 10^{20}$	$1.1 \cdot 10^{21}$
ConvS2S Ensemble [9]	26.36	41.29	$7.7 \cdot 10^{19}$	$1.2 \cdot 10^{21}$
Transformer (base model)	27.3	38.1	$3.3 \cdot 10^{18}$	
Transformer (big)	28.4	41.8	$2.3 \cdot 10^{19}$	

Thank you!

Q?