Machine Learning

Lecture 20: Convolution Neural Network

Chandresh Kumar Maurya
Assistant Professor in the CSE department

Indian Institute of Technology Indore, India

Slides Credit

Some of the slides are adapted from

- Dr. Piyush Rai machine learning course IITK CS 771A.
- Sudeep Raja, Columbia University
- Quoc Vi. Lee, A Tutorial on Deep Learning, Google Inc.

Table of Contents

Introduction

2 CNN

Limitations of Feed-Forward Networks

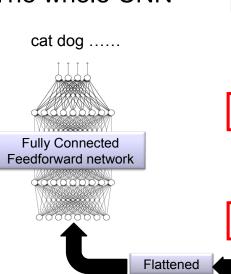
- Require a huge number of parameters (note that the consecutive layers are fully connected)
- Not ideal for data that exhibit locality structure, e.g., (e.g., images, sentences)
 - → Kind of works but would be better to exploit locality in the data more explicitly
- Doesn't have a "memory", so not ideal when modeling sequence of observations

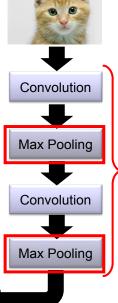
Table of Contents

Introduction

2 CNN

The whole CNN





Can repeat many times

CNN

- A feedforward neural network with a special structure
- Not all pairs of nodes are connected
- Weights are also "tied" (many connections have the same weights; color-coded above)
- The set of distinct weights defines a "filter" or "local" feature detector

CNN

- Applies 2 operations, convolution and pooling (subsampling), repeatedly on the input data
- ullet Convolution: Extract "local" properties of the signal. Uses a set of "filters" that have to be learned (these are the "weighted" W between layers)
- Pooling: Downsamples the outputs to reduce the size of representation
- Note: A nonlinearity is also introduced after the convolution layer

- An operation that captures local (e.g., spatial) properties of a signal
- Mathematically, the operation is defined as

$$h_{ij}^k = g((W^k * \boldsymbol{X}_{ij}) + b_k)$$

where W^k is the filter, * is the convolution operator, and g is nonlinearity

- ullet Usually several filters $\{W^k\}_{k=1}^K$ are applied (each will produce a separate "feature map").
- These filters have to be learned (these are the weights of the NN)

1		0	0	0	0	1
0)	1	0	0	1	0
0		0	1	1	0	0
1		0	0	0	1	0
0)	1	0	0	1	0
0		0	1	0	1	0

6 x 6 image

These are the network parameters to be learned.

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1



Filter 2

: :

Each filter detects a small pattern (3 x 3).

1	-1	-1	
-1	1	-1	
-1	-1	1	

Filter 1

stride=1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	0	0	0	1	0

Dot product 3

-1

6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

If stride=2

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
				l	

3 -3

6 x 6 image



Filter 1

stride=1

1	V	0	0	0	0	1
	0	X	0	0	1	0
	0	0		1	0	0
	V	0	0	0	1	0
	0	1	0	0	1	0
	0	0	1	0	1	0

6 x 6 image





Pooling/Downsampling

- Used to "downsample" the representation-size after convolution step.
- Also ensures robustness against minor rotations, shifts, corruptions in the image
- Popular approaches: Max-pooling, averaging pooling, etc

Max Pooling

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1



Filter 2

3 (-1)	-3 (-1
-3 1	0 -3
-3 -3	0 1

-1 -1 -1 -2 1



Why Pooling

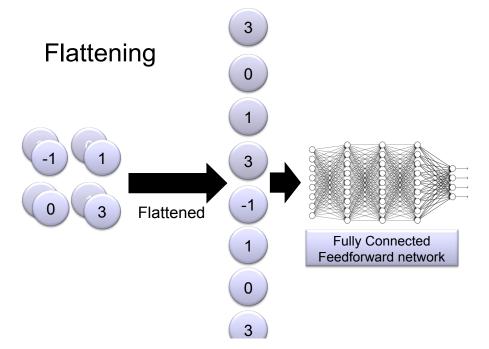
 Subsampling pixels will not change the object bird



We can subsample the pixels to make image smaller fewer parameters to characterize the image

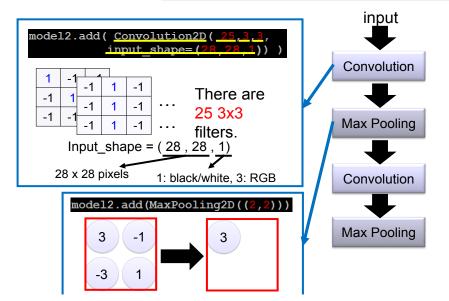
Strides

- Stride defines the number of nodes a filter moves between two consecutive convolution operations
- Likewise, we have a stride to define the same when applying pooling



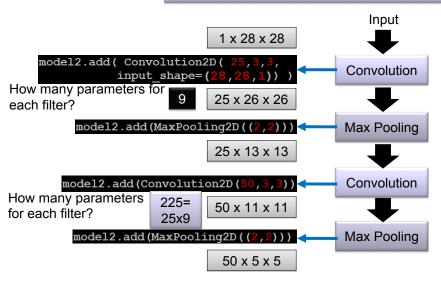
CNN in Keras

Only modified the *network structure* and *input format (vector -> 3-D tensor)*



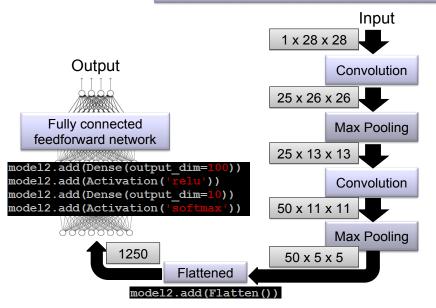
CNN in Keras

Only modified the *network structure* and *input format (vector -> 3-D array)*



CNN in Keras

Only modified the **network structure** and **input format (vector -> 3-D array)**



CNN for Sentiment Classification

See demo folder in PPT/Talks/DL

Bibliography I