

Recommender Systems for Biomedical and Health informatics

Presented by:
Chandresh Kumar Maurya
Assistant Professor at IIT Indore

Outline

- Can we improve medicine?
- What does this have to do with recsys?
- Medical decision support systems
- ML medical decision systems



Medicine: what can be improved?

Medical decisions require knowledge & data

- Doctors have partial information
 - Of the patient's history
 - Of the patient's symptoms
 - Of medical knowledge
 - Different demographics
 - Latest research findings
 - Not easy to remember all past information
- They also might have cognitive biases
 - They decide mostly based on past experience, but this experience is very limited
- Incentives of medical industry also play a role

Experts disagree

- Experts disagree among themselves.
 - E.g. oncology experts disagree on the value of colon screening
- Things treated as facts for years end up being wrong.
 - E.g. giving aspirin to reduce fever has been proved to be dangerous
- Psychiatric disorder diagnosis even lower agreement
(Cohen's Kappa of 0.2 or 0.3 in most cases)

ORIGINAL ARTICLE



A Decade of Reversal: An Analysis of 146 Contradicted Medical Practices

Vinay Prasad, MD; Andrae Vandross, MD; Caitlin Toomey, MD; Michael Cheung, MD; Jason Rho, MD; Steven Quinn, MD; Satish Jacob Chacko, MD; Durga Borkar, MD; Victor Gall, MD; Senthil Selvaraj, MD; Nancy Ho, MD; and Adam Cifu, MD


Abstract

Objective: To identify medical practices that offer no net benefits.

Methods: We reviewed all original articles published in 10 years (2001-2010) in one high-impact journal. Articles were classified on the basis of whether they addressed a medical practice, whether they tested a new or existing therapy, and whether results were positive or negative. Articles were then classified as 1 of 4 types: replacement, when a new practice surpasses standard of care; back to the drawing board, when a new practice is no better than current practice; reaffirmation, when an existing practice is found to be better than a lesser standard; and reversal, when an existing practice is found to be no better than a lesser therapy. This study was conducted from August 1, 2011, through October 31, 2012.

Results: We reviewed 2044 original articles, 1344 of which concerned a medical practice. Of these, 881 articles (73.0%) examined a new medical practice, whereas 363 (27.0%) tested an established practice. A total of 947 studies (70.5%) had positive findings, whereas 397 (29.5%) reached a negative conclusion. A total of 756 articles addressed a medical practice constituted replacement. 165 were back to the drawing

the value of second opinions



Growing complexity of knowledge

- Very slow pace for doctors to acquire new knowledge
- Increased amount of complexity
 - E.g. Coordinating specialists treating a patient is very complex and done using very little technology
- No universally accepted processes and recommendations to make decisions.
 - Those depend on each institution and practitioner
 - 50% of the recommendations made in guidelines based on expert opinion, case studies, or standards of care, not systematic studies

No personalization!

- Clinical Practice Guides are not personalized
 - They might prove negative since they fail to take into account interactions between different diseases (e.g. in older patients)
- Research done on “homogeneous”, healthy subjects
- It is very hard for doctors to “manually” personalize their “recommendations”



What does this have to do with Recsys?



A word on Precision Medicine

Precision medicine

- According to the National Institutes of Health (NIH), precision medicine is:

"an emerging approach for disease treatment and prevention that takes into account individual variability in genes, environment, and lifestyle for each person."



The NEW ENGLAND JOURNAL of MEDICINE

Perspective
FEBRUARY 26, 2015

A New Initiative on Precision Medicine

Francis S. Collins, M.D., Ph.D., and Harold Varmus, M.D.

Tonight, I'm launching a new Precision Medicine Initiative to bring us closer to curing diseases like cancer and diabetes — and to give all of us access to the personalized information we need to keep ourselves and our families healthier."

— President Barack Obama, State of the Union Address, January 20, 2015

President Obama has long expressed a strong conviction that science offers great potential for improving health. Now, the President has announced a research initiative that aims to accelerate prog-

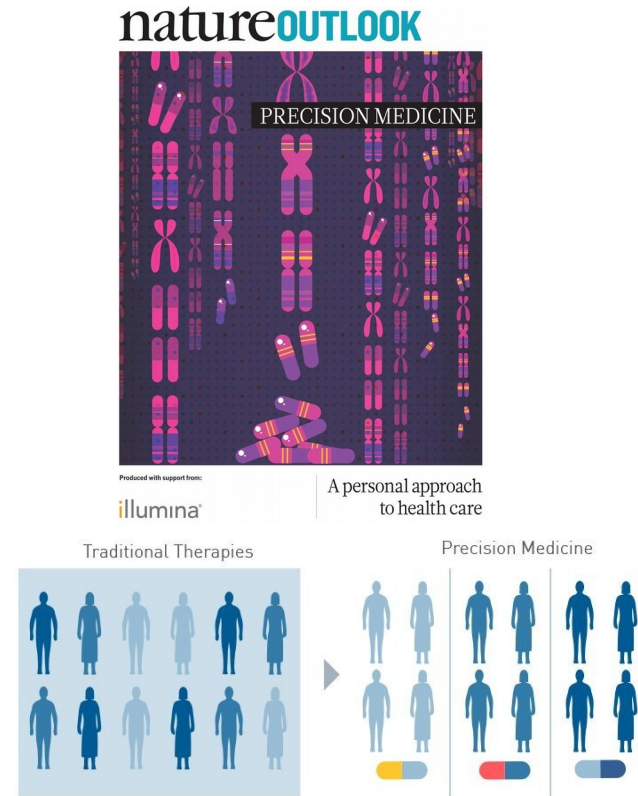
variability into account — is not new; blood typing, for instance, has been used to guide blood transfusions for more than a century. But the prospect of applying this concept broadly has been

is a broad research program to encourage creative approaches to precision medicine, test them rigorously, and ultimately use them to build the evidence base needed to guide clinical practice.

The proposed initiative has two main components: a near-term focus on cancers and a longer-term aim to generate knowledge applicable to the whole range of health and disease. Both components are now within our reach because of advances in basic research, including molecular biol-

Precision medicine

- Goal: predict more accurately treatment and prevention given a particular disease & group of people
- In contrast to "one-size-fits-all" approach where treatment & prevention are developed for average person
- Term is relatively new, but concept has been around for many years.
 - *E.g. blood transfusion is not given from a randomly selected donor*



Recsys & Medicine

Recommender System (Recsys)

- Recommender systems aim to help users by providing suitable options to execute a task easily and efficiently.
- Such systems learn user behavior by filtering through a large amount of data

Recsys in Health Care

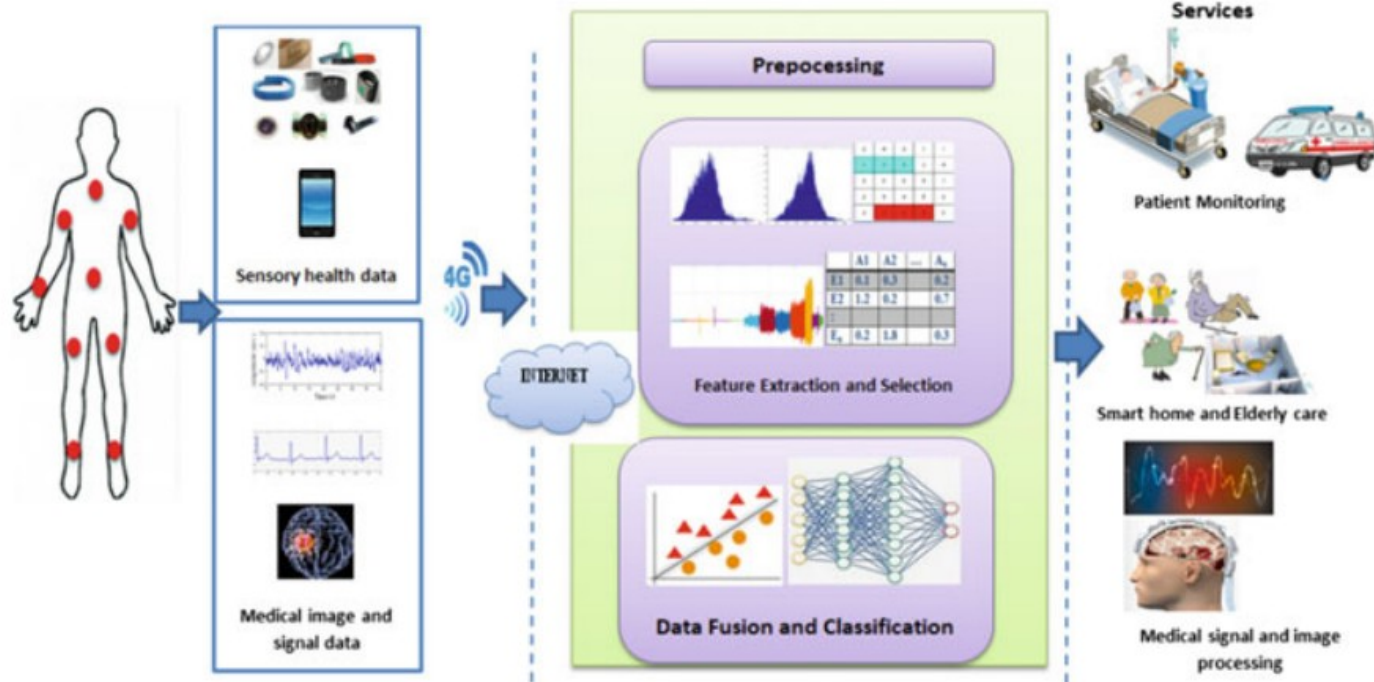


Fig. 2 System architecture of health recommender system

Table 1 Overview of the applications of recommendations for health informatics

Area of applications	Application	Input parameters	Learning techniques
Bioinformatics [4–6]	Drug design	Molecule compounds	Deep neural network
	RNA binding protein Compound protein interaction	Gene RNA/DNA sequences Molecule compounds	Deep belief network
			Deep neural network
Medical imaging [7–9]	Tissue classification Organ segmentation Tumor detection Hemorrhage detection	MRI/CT images Microscopy Hyperspectral images Endoscopy images	Convolutional neural network
			Convolutional deep belief network
			Deep neural network
Pervasive sensing [10–13]	Monitoring of biological parameters Anomaly detection	ECG, EEG Devices implanted	Convolutional neural network
	Human activity recognition	Wearable sensing devices Smartphones Video	Convolutional neural network
			Deep belief network
	Obstacle detection Sign language recognition Hand gesture recognition	RGB-D camera Real-sense camera Depth camera	Convolutional neural network
			Deep belief network
Public health [14, 15]	Lifestyle diseases Infectious disease epidemics	Text messages Social media data Geo-tagged images	Convolutional neural network
			Deep belief network
			Deep neural network

How to improve medicine?

- More/better personalization
- Make better use of (big) data
- Automate processes through algorithms and machine learning
- Better user experience both for patients and doctors

Does this sound familiar?

Int. J. Environ. Res. Public Health **2014**, *11*, 2580–2607; doi:10.3390/ijerph110302580

OPEN ACCESS

International Journal of
Environmental Research and
Public Health
ISSN 1660-4601
www.mdpi.com/journal/ijerph

Article

Health Recommender Systems: Concepts, Requirements, Technical Basics and Challenges

Martin Wiesner * and Daniel Pfeifer

Department of Medical Informatics, Heilbronn University, Max-Planck-Str. 39, Heilbronn 74081, Germany; E-Mail: daniel.pfeifer@hs-heilbronn.de

* Author to whom correspondence should be addressed; E-Mail: martin.wiesner@hs-heilbronn.de; Tel.: +49-7131-504-6947.

Received: 3 December 2013; in revised form: 4 February 2014 / Accepted: 8 February 2014 / Published: 3 March 2014


Abstract: During the last decades huge amounts of data have been collected in clinical databases representing patients' health states (e.g., as laboratory results, treatment plans, medical reports). Hence, digital information available for patient-oriented decision making has increased drastically but is often scattered across different sites. As a solution, *personal health record systems* (PHRS) are meant to centralize an individual's health data and to allow access for the owner as well as for authorized health professionals. Yet, expert-oriented language, complex interrelations of medical facts and information overload in general pose major obstacles for patients to understand their own record and to draw adequate conclusions. In this context, *recommender systems* may supply patients with additional laymen-friendly information helping to better comprehend their health status as represented by their record.



Medical Decision Support Systems

Medical Diagnosis

- Diagnosis:
 - *“a mapping from a patient’s data (normal and abnormal history , physical examination , and laboratory data) to a nosology of disease states the process of determining by examination the nature and circumstances of a diseased condition”*
 - *“The knowledge of how to “work up” the patient depends critically on the ability to evoke history, symptoms, and physical examination findings, concurrently with the ability to generate diagnostic hypotheses that suggest how to further refine or pursue the findings already elicited, or to pursue completely different additional findings. In addition, this must be done in a compassionate and cost-effective manner” (R.A. Miller 1990)*



Diagnosis Decision Support Systems

- DDSS

- A computer - based algorithm that assists a clinician with one or more component steps of the diagnostic process
- It involves diverse activities:
 - information gathering
 - pattern recognition
 - problem solving
 - Decision-making
 - judgment under uncertainty
 - Empathy
- Large amounts of highly organized knowledge are necessary

- Two categories

- General-domain DDSS
- Specialized DDSS

Medical Diagnosis as a Recsys


- Input signals
 - Implicit: Data coming from sensors, labs...
 - Explicit: What does the patient tell us
 - Others: Demographic, family history...
- Algorithms
 - Find what worked for “similar” patients in a “similar” situation
- Output
 - Ranked list with “likely” options
 - Need to also worry about other variables such as acuity, explanations....



Knowledge-based approaches to medical decision systems

Decision/Diagnosis support systems

- They have been developed for decades
- Many early DDSS based on Bayesian reasoning (60s-70s)
 - Bayesian networks (80s-90s)
 - Neural networks (lately)
- Most of them have been developed “manually” with doctors looking at research and manually encoding that into knowledge bases.





Data for Medical Decision Systems

Knowledge bases

- Knowledge base construction
 - Build model from existing medical knowledge
 - E.g. using doctors to read on medical literature and encode “well-established facts”
- Data = Medical Research
 - Ingesting and understanding medical publications can be mostly automated



Electronic Health Records

- EHR/EMRs include digital information about patients *encounters* with doctors or the health system

The screenshot displays a 'Patient Chart' window for a user named 'Demo.F. alther' (ID: 5465) on 05-Mar-1995. The chart is titled 'CHART REVIEW' and '28-Jul-2005 10:15'. The 'Problem List' is active, showing two entries: 'TYPE 2 DIABETES MELLITUS' (Status: Active, Entered: 03/11/2000, Onset: 03/11/2000) and 'HYPERTENSION' (Status: Active, Entered: 02/04/2000, Onset: 01/19/1999). Below the problem list are 'ICD Pick Lists' for Administrative, Medicine, and Pediatrics. The 'Historical Diagnosis' section shows a list of past encounters with dates, ICD codes, and descriptions. The bottom of the window features a navigation bar with tabs for 'Notifications', 'Cover Sheet', 'Triage', 'Wellness', 'Notes', 'Prob/POV', 'Orders', 'Medications', 'Labs', 'D/C Summ', 'Reports', and 'Consults'. The 'Prob/POV' tab is currently selected.

128

© 2008

IMIA and Schottauer GmbH

Extracting Information from Textual Documents in the Electronic Health Record: A Review of Recent Research

S. M. Meystre¹, G. K. Savova², K. C. Kipper-Schuler², J. F. Hurdle¹

¹ Department of Biomedical Informatics, University of Utah School of Medicine, Salt Lake City, Utah, USA

² Biomedical Informatics Research, Mayo Clinic College of Medicine, Rochester, Minnesota, USA

Summary

Objectives: We examine recent published research on the extraction of information from textual documents in the Electronic Health Record (EHR).

Methods: Literature review of the research published after 1995, based on PubMed and conference proceedings.

Introduction

In the biomedical domain, the rapid adoption of Electronic Health Records (EHR) with the parallel growth of narrative data in electronic form, alone

rules or based on statistical methods and machine learning. The information extracted can then be linked to concepts in standard terminologies and used for coding. The information can also be used for decision support and to enrich

Generating Multi-label Discrete Patient Records using Generative Adversarial Networks

Edward Choi¹

MP2893@GATECH.EDU

Siddharth Biswal¹

SBISWAL7@GATECH.EDU

Bradley Malin²

BRADLEY.MALIN@VANDERBILT.EDU

Jon Duke¹

JON.DUKE@GATECH.EDU

Walter F. Stewart³

STEWARTWF@SUTTERHEALTH.ORG

Jimeng Sun¹

JISUN@CC.GATECH.EDU

¹ GEORGIA INSTITUTE OF TECHNOLOGY

² VANDERBILT UNIVERSITY

³ SUTTER HEALTH

Abstract

Access to electronic health record (EHR) data has motivated computational advances in medical research. However, various concerns, particularly over privacy, can limit access to and collaborative use of EHR data. Sharing synthetic EHR data could mitigate risk.

In this paper, we propose a new approach, medical Generative Adversarial Network (medGAN), to generate realistic synthetic patient records. Based on input real patient records, medGAN can generate high-dimensional discrete variables (e.g., binary and count features) via a combination of an autoencoder and generative adversarial networks. We also propose a method to generate synthetic EHR data that can be used for research purposes.

Ontologies

- Snomed Clinical Terms
 - Computer processable collection of medical terms providing codes, terms, synonyms and definitions used in clinical documentation and reporting.
 - Considered to be the most comprehensive, multilingual clinical healthcare terminology
 - Primary purpose: encode the meanings that are used in health information & support effective clinical recording of data with the aim of improving patient care.
 - It provides the core general terminology for electronic health records.
 - It includes: clinical findings, symptoms, diagnoses, procedures, body structures, organisms and other etiologies, substances, pharmaceuticals, devices and specimens.



Parent(s): (Select a parent to make it the "Current Concept".) Viral upper respiratory tract infection (disorder)	Current Concept: Fully Specified Name: Common cold (disorder) ConceptId: 82272006
Current Concept: Common cold (disorder)	Defining Relationships: Is a Viral upper respiratory tract infection (disorder) Causative agent Virus (organism) Finding site Upper respiratory tract structure (body structure) Pathological process Infectious process (qualifier value) <small>This concept is primitive.</small>
Child(ren): (N=0) (Select a child to make it the "Current Concept".)	Qualifiers: View Qualifying Characteristics and Facts
	Descriptions (Synonyms): Fully Specified Name: Common cold (disorder) Preferred: Common cold Synonym: Acute coryza Synonym: Acute nasal catarrh Synonym: Acute rhinitis Synonym: Infective rhinitis Synonym: Acute nasopharyngitis Synonym: Infective nasopharyngitis Synonym: Head cold Synonym: Acute infective rhinitis Synonym: Cold Synonym: Acute nasopharyngitis, NOS Synonym: Infective nasopharyngitis, NOS

Ontologies

- ICD-10

- 10th revision of the International Statistical Classification of Diseases and Related Health Problems (ICD), a medical classification list by the World Health Organization (WHO)
- It contains codes for diseases, signs and symptoms, abnormal findings, complaints, social circumstances, and external causes of injury or diseases
- The code set allows more than 14,400 different codes and permits the tracking of new diagnoses.



Ontologies

- UMLS (1986)
 - Designed and is maintained by the US National Library of Medicine, updated quarterly, free
 - Compendium of many controlled vocabularies in the biomedical sciences.
 - Provides mapping structure among vocabularies
 - Allows to translate among the various terminology systems
 - It can be interpreted as a thesaurus and ontology of biomedical concepts.
 - It provides tools for natural language processing.
 - Intended for developers of systems in medical informatics.



Unified Medical Language System® (UMLS®)

Combining and aggregating data

- At the end, what you need is a way to process and aggregate all these different sources of data either manually (expert systems) or algorithmically (ML)



01
02
03
04
05
06
07
08
09
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131
132
133
134
135
136
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160
161
162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208
209
210
211
212
213
214
215
216
217
218
219
220
221
222
223
224
225
226
227
228
229
230
231
232
233
234
235
236
237
238
239
240
241
242
243
244
245
246
247
248
249
250
251
252
253
254
255
256
257
258
259
260
261
262
263
264
265
266
267
268
269
270
271
272
273
274
275
276
277
278
279
280
281
282
283
284
285
286
287
288
289
290
291
292
293
294
295
296
297
298
299
300
301
302
303
304
305
306
307
308
309
310
311
312
313
314
315
316
317
318
319
320
321
322
323
324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341
342
343
344
345
346
347
348
349
350
351
352
353
354
355
356
357
358
359
360
361
362
363
364
365
366
367
368
369
370
371
372
373
374
375
376
377
378
379
380
381
382
383
384
385
386
387
388
389
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407
408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431
432
433
434
435
436
437
438
439
440
441
442
443
444
445
446
447
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464
465
466
467
468
469
470
471
472
473
474
475
476
477
478
479
480
481
482
483
484
485
486
487
488
489
490
491
492
493
494
495
496
497
498
499
500
501
502
503
504
505
506
507
508
509
510
511
512
513
514
515
516
517
518
519
520
521
522
523
524
525
526
527
528
529
530
531
532
533
534
535
536
537
538
539
540
541
542
543
544
545
546
547
548
549
550
551
552
553
554
555
556
557
558
559
560
561
562
563
564
565
566
567
568
569
570
571
572
573
574
575
576
577
578
579
580
581
582
583
584
585
586
587
588
589
590
591
592
593
594
595
596
597
598
599
600
601
602
603
604
605
606
607
608
609
610
611
612
613
614
615
616
617
618
619
620
621
622
623
624
625
626
627
628
629
630
631
632
633
634
635
636
637
638
639
640
641
642
643
644
645
646
647
648
649
650
651
652
653
654
655
656
657
658
659
660
661
662
663
664
665
666
667
668
669
670
671
672
673
674
675
676
677
678
679
680
681
682
683
684
685
686
687
688
689
690
691
692
693
694
695
696
697
698
699
700
701
702
703
704
705
706
707
708
709
710
711
712
713
714
715
716
717
718
719
720
721
722
723
724
725
726
727
728
729
730
731
732
733
734
735
736
737
738
739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754
755
756
757
758
759
760
761
762
763
764
765
766
767
768
769
770
771
772
773
774
775
776
777
778
779
780
781
782
783
784
785
786
787
788
789
790
791
792
793
794
795
796
797
798
799
800
801
802
803
804
805
806
807
808
809
810
811
812
813
814
815
816
817
818
819
820
821
822
823
824
825
826
827
828
829
830
831
832
833
834
835
836
837
838
839
840
841
842
843
844
845
846
847
848
849
850
851
852
853
854
855
856
857
858
859
860
861
862
863
864
865
866
867
868
869
870
871
872
873
874
875
876
877
878
879
880
881
882
883
884
885
886
887
888
889
890
891
892
893
894
895
896
897
898
899
900
901
902
903
904
905
906
907
908
909
910
911
912
913
914
915
916
917
918
919
920
921
922
923
924
925
926
927
928
929
930
931
932
933
934
935
936
937
938
939
940
941
942
943
944
945
946
947
948
949
950
951
952
953
954
955
956
957
958
959
960
961
962
963
964
965
966
967
968
969
970
971
972
973
974
975
976
977
978
979
980
981
982
983
984
985
986
987
988
989
990
991
992
993
994
995
996
997
998
999
1000



ML for Medical Decision Systems

Health graphs

SCIENTIFIC DATA

OPEN

SUBJECT CATEGORIES

- » Data mining
- » Diagnosis
- » Epidemiology
- » Outcomes research

Building the graph of medicine from millions of clinical narratives

Samuel G. Finlayson¹, Paea LePendou¹ & Nigam H. Shah

Electronic health records (EHR) represent a rich and relatively untapped resource for characterizing the true nature of clinical practice and for quantifying the degree of inter-relatedness of medical entities such as

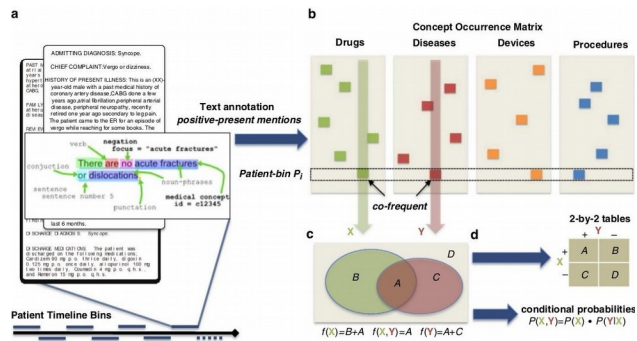


Figure 1. Workflow Architecture. The architecture of our workflow starts with (a) patient notes that are grouped together based on their nearness in time. Given the patient timeline bins, clinical terms are recognized from the notes and recorded into (b) the clinical concept occurrence matrix, which is scanned for (c) counting pairwise the frequency and co-frequency of concepts. This data can be used to calculate (d) contingency tables and Bayesian probability estimates. For example, the concept X has a frequency of $f(X)$ and is pairwise co-frequent with concept Y exactly $f(X,Y)$ times.

SCIENTIFIC REPORTS

OPEN

Learning a Health Knowledge Graph from Electronic Medical Records

Maya Rotmenski¹, Yoni Halpern¹, Abdulhakim Tilmat¹, Steven Hong^{1,4} & David Sontag^{1,5,6}

Demand for clinical decision support systems in medicine and self-diagnostic symptom checkers has substantially increased in recent years. Existing platforms rely on knowledge bases manually compiled through a labor-intensive process or automatically derived using simple pairwise statistics. This study explored an automated process to learn high quality knowledge bases linking diseases and symptoms directly from electronic medical records. Medical concepts were extracted from 273,174 de-identified patient records and maximum likelihood estimation of three probabilistic models was used to automatically construct knowledge graphs: logistic regression, naive Bayes classifier and a Bayesian network using noisy OR gates. A graph of disease-symptom relationships was elicited from the learned parameters and the constructed knowledge graphs were evaluated and validated, with permission, against Google's manually-constructed knowledge graph and against expert physician opinions. Our study shows that direct and automated construction of high quality health knowledge graphs from medical records using rudimentary concept extraction is feasible. The noisy OR model produces a high quality knowledge graph reaching precision of 0.85 for a recall of 0.6 in the clinical evaluation. Noisy OR significantly outperforms all tested models across evaluation frameworks ($p < 0.01$).

Received: 3 March 2017
Accepted: 1 June 2017
Published online: 20 July 2017

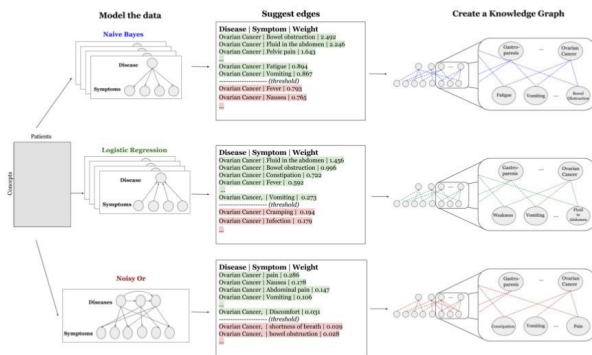


Figure 2. Workflow of modeling the relationship between diseases and symptoms and knowledge graph construction, for each of our 3 models (naive Bayes, logistic regression and noisy OR).

- Understanding what doctors say
- Understanding what patients say

Diagnostic Inferencing via Improving Clinical Concept Extraction with Deep Reinforcement Learning: A Preliminary Study

Yuan Ling
Sadid A. Hasan
Vivek Datla
Ashegul Qadir
Kathy Lee
Joey Liu
Oladimeji Farri

Artificial Intelligence Laboratory, Philips Research North America
Cambridge, MA, USA

YUAN.LING@PHILIPS.COM
SADID.HASAN@PHILIPS.COM
VIVEK.DATLA@PHILIPS.COM
ASHEGUL.QADIR@PHILIPS.COM
KATHY.LEE_1@PHILIPS.COM
JOEY.LIU@PHILIPS.COM
DIMEJI.FARRI@PHILIPS.COM




Figure 2. Mappings among terms and concepts. The figure explains the mappings that can be used to decode the frequency files stored in records 1 and 2. We use a subset of terms related to 'hydrocephalus' to demonstrate the mapping of terms (File 1) to concepts and UMLS CUIs. Terms map onto concepts in a many-to-many fashion (File 1). Concepts map onto CUIs in a one-to-one fashion (File 2b) and have an associated string for human readability (File 2a).

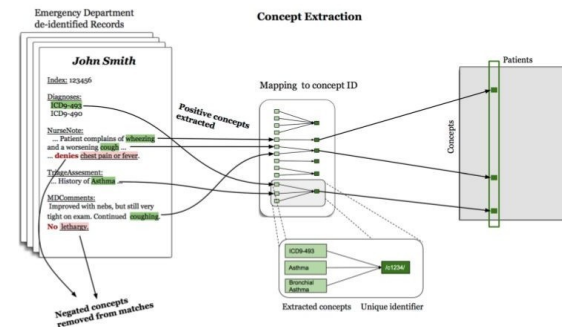


Figure 1. Concept extraction pipeline. Non-negated concepts and ICD-9 diagnosis codes are extracted from Emergency Department electronic medical records. Concepts, codes and concept aliases are mapped to unique IDs, which in turn populate a co-occurrence matrix of size (Concepts) \times (Patients).

Diagnostic Systems

● Building diagnostic systems from data through ML

Proceedings of Machine Learning for Healthcare 2016

Doctor AI: Predicting Clinical via Recurrent Neural Networks

Edward Choi, Mohammad Taha Bahadori
College of Computing
Georgia Institute of Technology
Atlanta, GA, USA

Andy Schuetz, Walter F. Stewart
Research Development & Dissemination
Sutter Health
Walnut Creek, CA, USA

Jimeng Sun
College of Computing
Georgia Institute of Technology
Atlanta, GA, USA

SCHU

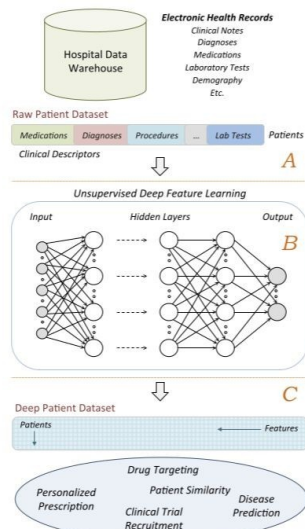


Figure 1. Conceptual framework used to derive the deep patient representation through unsupervised deep learning of a large EHR data warehouse. (A) Pre-processing stage to obtain raw patient representations from the EHRs. (B) The raw representations are modeled by the unsupervised deep architecture leading to a set of general and robust features. (C) The deep features are applied to the entire hospital database to derive patient representations that can be applied to a number of clinical tasks.

SCIENTIFIC REPORTS

OPEN

Deep Patient: An Unsupervised Representation to Predict the Future of Patients from the Electronic Health Records

Received: 28 January 2016

Accepted: 27 April 2016

Published: 17 May 2016

Riccardo Miotto^{1,2,3}, Li Li^{1,2,3}, Brian A. Kilday^{1,2,3}, Joel T. Dudley^{1,2,3}

Secondary use of electronic health records (EHRs) promises to advance clinical research and better inform clinical decision making. Challenges in summarizing and representing patient data prevent widespread practice of predictive modeling using EHRs. Here we present a novel unsupervised deep feature learning method to derive a general-purpose patient representation from EHR data that facilitates clinical predictive modeling. In particular, a three-layer stack of denoising autoencoders was used to capture hierarchical regularities and dependencies in the aggregated EHRs of about 700,000 patients from the Mount Sinai data warehouse. The result is a representation we name "deep patient". We evaluated this representation as broadly predictive of health states by assessing the probability of patients to develop various diseases. We performed evaluation using 76,214 test patients comprising 78 diseases from diverse clinical domains and temporal windows. Our results significantly outperformed those achieved using representations based on raw EHR data and alternative feature learning strategies. Prediction performance for severe diabetes, schizophrenia, and various cancers were among the top performing. These findings indicate that deep learning applied to EHRs can derive patient representations that offer improved clinical predictions, and could provide a machine learning framework for augmenting clinical decision systems.

Challenges

- Algorithmic: e.g. combining expert rule-based and ML
- Data: quality, sparsity, and bias in data
- UX: trustworthiness and engagement of the system, incentives...
- Legal
- ...

It's about time we overcome all of these.

Conclusions

- Medicine/healthcare is one of the areas in society that can benefit most from technology in general and algorithmic approaches in particular
- Medicine has many similarities to recommender systems
 - Importance of data (both big and small)
 - Content and collaborative approaches
 - Importance of UI/UX (e.g. user understanding, feedback, and explanations)
 - Evaluation has to be sensitive to overall optimization problem
- Huge opportunity, great mission: looking forward to many advances that will save millions of lives in the next few years

References

- *“Probabilistic diagnosis using a reformulation of the INTERNIST-1/QMR knowledge base”* . Shwe et al. 1991.
- *“Computer-assisted diagnostic decision support: history, challenges, and possible paths forward”* Miller. 2009.
- *“Mining Biomedical Ontologies and Data Using RDF Hypergraphs”* Liu et al. 2013.
- *“Health Recommender Systems: Concepts, Requirements, Technical Basics & Challenges”*, Wiesner & Pfeifer, 2014.
- *“A ‘Green Button’ For Using Aggregate Patient Data At The Point Of Care ”* Longhurst et al. 2014.
- *“Building the graph of medicine from millions of clinical narratives”* Finlayson et al. 2014.
- *“Comparison of Physician and Computer Diagnostic Accuracy”* Semigran et al. 2016.
- *“Identifiable Phenotyping using Constrained Non-Negative Matrix Factorization”*. Joshi et al. 2016.
- *“Clinical Tagging with Joint Probabilistic Models”* . Halpern et al. 2016.
- *“Deep Patient: An Unsupervised Representation to Predict the Future of Patients from EHR ”*. Miotto et al. 2016.
- *“Learning a Health Knowledge Graph from Electronic Medical Records”* Rotmensch et al. 2017.
- *“Clustering Patients with Tensor Decomposition”*. Ruffini et al. 2017.
- *“Patient Similarity Using Population Statistics and Multiple Kernel Learning ”*. Conroy et al. 2017.
- *“Diagnostic Inferencing via Clinical Concept Extraction with Deep Reinforcement Learning”*. Ling et al. 2017.
- *“Generating Multi-label Discrete Patient Records using Generative Adversarial Networks”* Choi et al. 2017

Saha J., Chowdhury C., Biswas S. (2020) Review of Machine Learning and Deep Learning Based Recommender Systems for Health Informatics. In: Dash S., Acharya B., Mittal M., Abraham A., Kelemen A. (eds) Deep Learning Techniques for Biomedical and Health Informatics. Studies in Big Data, vol 68. Springer, Cham. https://doi.org/10.1007/978-3-030-33966-1_6